



# 1. Business Problem

## 1.1 Problem Description

Netflix is all about connecting people to the movies they love. To help customers find those movies, the recommendation system: CinematchSM. Its job is to predict whether someone will enjoy a movie based on other movies. Netflix uses those predictions to make personal movie recommendations based on each. **Cinematch** is doing pretty well, it can always be made better.

Now there are a lot of interesting alternative approaches to how Cinematch works that Netflix haven't tried some aren't. We're curious whether any of these can beat Cinematch by making better predictions. Because approach it could make a big difference to our customers and our business.

Credits: <https://www.netflixprize.com/rules.html>

## 1.2 Problem Statement

Netflix provided a lot of anonymous rating data, and a prediction accuracy bar that is 10% better than training data set. (Accuracy is a measurement of how closely predicted ratings of movies match user's ratings)

## 1.3 Sources

- <https://www.netflixprize.com/rules.html>
- <https://www.kaggle.com/netflix-inc/netflix-prize-data>
- Netflix blog: <https://medium.com/netflix-techblog/netflix-recommendations-beyond-the-5-stars-4c9c63610590>
- surprise library: <http://surpriselib.com/> (we use many models from this library)
- surprise library doc: [http://surprise.readthedocs.io/en/stable/getting\\_started.html](http://surprise.readthedocs.io/en/stable/getting_started.html) (we use many models from this library)
- installing surprise: <https://github.com/NicolasHug/Surprise#installation>
- Research paper: <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf> (most of the models in surprise are based on this paper)
- SVD Decomposition : <https://www.youtube.com/watch?v=P5mlg91as1c>

## 1.4 Real world/Business Objectives and constraints

## Objectives:

1. Predict the rating that a user would give to a movie that he has not yet rated.
2. Minimize the difference between predicted and actual rating (RMSE and MAPE)

## Constraints:

1. Some form of interpretability.

# 2. Machine Learning Problem

## 2.1 Data

### 2.1.1 Data Overview

Get the data from : <https://www.kaggle.com/netflix-inc/netflix-prize-data/data>

Data files :

- combined\_data\_1.txt
- combined\_data\_2.txt
- combined\_data\_3.txt
- combined\_data\_4.txt
- movie\_titles.csv

The first line of each file [combined\_data\_1.txt, combined\_data\_2.txt, combined\_data\_3.txt,

CustomerID,Rating,Date

MovieIDs range from 1 to 17770 sequentially.

CustomerIDs range from 1 to 2649429, with gaps. There are 480189 users.

Ratings are on a five star (integral) scale from 1 to 5.

Dates have the format YYYY-MM-DD.

### 2.1.2 Example Data point

1:  
1488844,3,2005-09-06  
822109,5,2005-05-13  
885013,4,2005-10-19  
30878,4,2005-12-26  
823519,3,2004-05-03  
893988,3,2005-11-17  
124105,4,2004-08-05  
1248029,3,2004-04-22  
1842128,4,2004-05-09  
2238063,3,2005-05-11  
1503895,4,2005-05-19  
2207774,5,2005-06-06  
2590061,3,2004-08-12  
2442,3,2004-04-14  
543865,4,2004-05-28  
1209119,4,2004-03-23  
804919,4,2004-06-10  
1086807,3,2004-12-28  
1711859,4,2005-05-08  
372233,5,2005-11-23  
1080361,3,2005-03-28  
1245640,3,2005-12-19  
558634,4,2004-12-14  
2165002,4,2004-04-06  
1181550,3,2004-02-01  
1227322,4,2004-02-06  
427928,4,2004-02-26  
814701,5,2005-09-29  
808731,4,2005-10-31  
662870,5,2005-08-24  
337541,5,2005-03-23  
786312,3,2004-11-16  
1133214,4,2004-03-07  
1537427,4,2004-03-29  
1209954,5,2005-05-09  
2381599,3,2005-09-12  
525356,2,2004-07-11  
1910569,4,2004-04-12  
2263586,4,2004-08-20  
2421815,2,2004-02-26  
1009622,1,2005-01-19

1481961,2,2005-05-24  
401047,4,2005-06-03  
2179073,3,2004-08-29  
1434636,3,2004-05-01  
93986,5,2005-10-06  
1308744,5,2005-10-29  
2647871,4,2005-12-30  
1905581,5,2005-08-16  
2508819,3,2004-05-18  
1578279,1,2005-05-19  
1159695,4,2005-02-15  
2588432,3,2005-03-31  
2423091,3,2005-09-12  
470232,4,2004-04-08  
2148699,2,2004-06-05  
1342007,3,2004-07-16  
466135,4,2004-07-13  
2472440,3,2005-08-13  
1283744,3,2004-04-17  
1927580,4,2004-11-08  
716874,5,2005-05-06  
4326,4,2005-10-29

## 2.2 Mapping the real world problem to a Machine Learning Problem

### 2.2.1 Type of Machine Learning Problem

For a given movie and user we need to predict the rating would be given by him/her to the movie  
The given problem is a Recommendation problem  
It can also be seen as a Regression problem

### 2.2.2 Performance metric

- Mean Absolute Percentage Error: [https://en.wikipedia.org/wiki/Mean\\_absolute\\_percentage\\_error](https://en.wikipedia.org/wiki/Mean_absolute_percentage_error)
- Root Mean Square Error: [https://en.wikipedia.org/wiki/Root-mean-square\\_deviation](https://en.wikipedia.org/wiki/Root-mean-square_deviation)

### 2.2.3 Machine Learning Objective and Constraints

1. Minimize RMSE.
2. Try to provide some interpretability.

```
# this is just to know how much time will it take to run this entire ipython notebook
from datetime import datetime
# globalstart = datetime.now()
import pandas as pd
import numpy as np
import matplotlib
matplotlib.use('nbagg')

import matplotlib.pyplot as plt
plt.rcParams.update({'figure.max_open_warning': 0})

import seaborn as sns
sns.set_style('whitegrid')
import os
from scipy import sparse
from scipy.sparse import csr_matrix

from sklearn.decomposition import TruncatedSVD
from sklearn.metrics.pairwise import cosine_similarity
import random
```

Double-click (or enter) to edit

## 3. Exploratory Data Analysis

### 3.1 Preprocessing

#### 3.1.1 Converting / Merging whole data to required format: u\_i, m\_j, r\_ij

```
start = datetime.now()
if not os.path.isfile('data.csv'):
    # Create a file 'data.csv' before reading it
    # Read all the files in netflix and store them in one big file('data.csv')
    # We're reading from each of the four files and appending each rating to a global file 'tr
    data = open('data.csv', mode='w')

    row = list()
    files=['combined_data_1.txt','combined_data_2.txt',
           'combined_data_3.txt', 'combined_data_4.txt']
```

```
for file in files:
    print("Reading ratings from {}...".format(file))
    with open(file) as f:
        for line in f:
            del row[:] # you don't have to do this.
            line = line.strip()
            if line.endswith(':'):
                # All below are ratings for this movie, until another movie appears.
                movie_id = line.replace(':', '')
            else:
                row = [x for x in line.split(',')]
                row.insert(0, movie_id)
                data.write(','.join(row))
                data.write('\n')
    print("Done.\n")
data.close()
print('Time taken :', datetime.now() - start)
```

 Reading ratings from combined\_data\_1.txt...
Done.

Reading ratings from combined\_data\_2.txt...
Done.

Reading ratings from combined\_data\_3.txt...
Done.

Reading ratings from combined\_data\_4.txt...
Done.

Time taken : 0:11:33.449895

Double-click (or enter) to edit

```
print("creating the dataframe from data.csv file..")
df = pd.read_csv('data.csv', sep=',',
                  names=['movie', 'user', 'rating', 'date'])
df.date = pd.to_datetime(df.date)
print('Done.\n')
```

```
# we are arranging the ratings according to time.
print('Sorting the dataframe by date..')
df.sort_values(by='date', inplace=True)
print('Done..')
```

 creating the dataframe from data.csv file...
Done.

Sorting the dataframe by date..
Done..

```
df.head()
```

|                 | movie | user   | rating | date       |
|-----------------|-------|--------|--------|------------|
| <b>56431994</b> | 10341 | 510180 | 4      | 1999-11-11 |
| <b>9056171</b>  | 1798  | 510180 | 5      | 1999-11-11 |
| <b>58698779</b> | 10774 | 510180 | 3      | 1999-11-11 |
| <b>48101611</b> | 8651  | 510180 | 2      | 1999-11-11 |
| <b>81893208</b> | 14660 | 510180 | 2      | 1999-11-11 |

```
df.describe()['rating']
```

 count 1.004805e+08  
 mean 3.604290e+00  
 std 1.085219e+00  
 min 1.000000e+00  
 25% 3.000000e+00  
 50% 4.000000e+00  
 75% 4.000000e+00  
 max 5.000000e+00  
 Name: rating, dtype: float64

### 3.1.2 Checking for NaN values

```
# just to make sure that all Nan containing rows are deleted..  

print("No of Nan values in our dataframe : ", sum(df.isnull().any()))
```

 No of Nan values in our dataframe : 0

Double-click (or enter) to edit

### 3.1.3 Removing Duplicates

```
dup_bool = df.duplicated(['movie','user','rating'])  

dups = sum(dup_bool) # by considering all columns..( including timestamp)  

print("There are {} duplicate rating entries in the data..".format(dups))
```

 There are 0 duplicate rating entries in the data..

Double-click (or enter) to edit

### 3.1.4 Basic Statistics (#Ratings, #Users, and #Movies)

```
print("Total data ")
print("-"*50)
print("\nTotal no of ratings :",df.shape[0])
print("Total No of Users   :", len(np.unique(df.user)))
print("Total No of movies  :", len(np.unique(df.movie)))
```

 Total data

```
-----  
Total no of ratings : 100480507  
Total No of Users   : 480189  
Total No of movies  : 17770
```

## 3.2 Spliting data into Train and Test(80:20)

```
if not os.path.isfile('train.csv'):
    # create the dataframe and store it in the disk for offline purposes..
    df.iloc[:int(df.shape[0]*0.80)].to_csv("train.csv", index=False)

if not os.path.isfile('test.csv'):
    # create the dataframe and store it in the disk for offline purposes..
    df.iloc[int(df.shape[0]*0.80):].to_csv("test.csv", index=False)

train_df = pd.read_csv("train.csv", parse_dates=['date'])
test_df = pd.read_csv("test.csv")
```

### 3.2.1 Basic Statistics in Train data (#Ratings, #Users, and #Movies)

```
# movies = train_df.movie.value_counts()
# users = train_df.user.value_counts()
print("Training data ")
print("-"*50)
print("\nTotal no of ratings :",train_df.shape[0])
print("Total No of Users   :", len(np.unique(train_df.user)))
print("Total No of movies  :", len(np.unique(train_df.movie)))
```

 Training data

```
-----  
Total no of ratings : 80384405  
Total No of Users   : 405041  
Total No of movies  : 17424
```

### 3.2.2 Basic Statistics in Test data (#Ratings, #Users, and #Movies)

```
print("Test data ")
print("-"*50)
```

```
print("\nTotal no of ratings :",test_df.shape[0])
print("Total No of Users   :", len(np.unique(test_df.user)))
print("Total No of movies  :", len(np.unique(test_df.movie)))
```



Test data

```
-----  
Total no of ratings : 20096102  
Total No of Users   : 349312  
Total No of movies  : 17757
```

## 3.3 Exploratory Data Analysis on Train data

Double-click (or enter) to edit

```
# method to make y-axis more readable
def human(num, units = 'M'):
    units = units.lower()
    num = float(num)
    if units == 'k':
        return str(num/10**3) + " K"
    elif units == 'm':
        return str(num/10**6) + " M"
    elif units == 'b':
        return str(num/10**9) + " B"
```

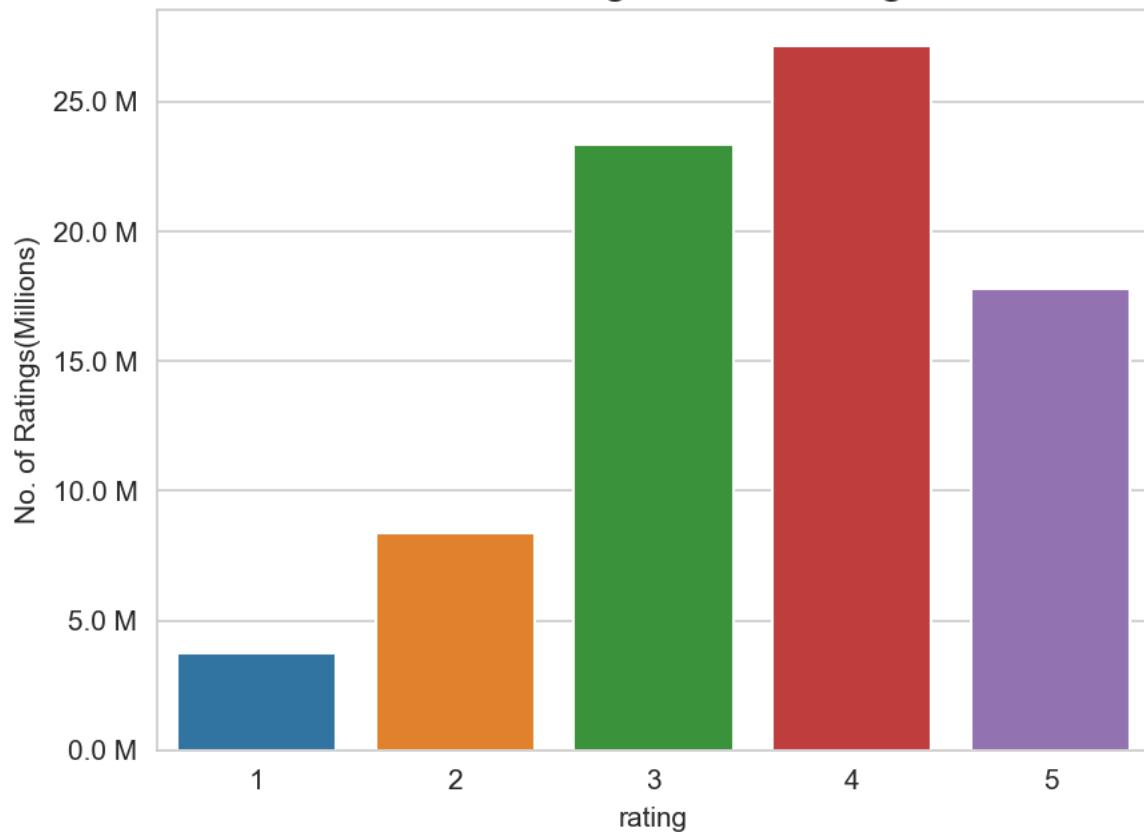
### 3.3.1 Distribution of ratings

```
fig, ax = plt.subplots()
plt.title('Distribution of ratings over Training dataset', fontsize=15)
sns.countplot(train_df.rating)
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
ax.set_ylabel('No. of Ratings(Millions)')

plt.show()
```



## Distribution of ratings over Training dataset



**Add new column (week day) to the data set for analysis.**

```
# It is used to skip the warning ''SettingWithCopyWarning''..
pd.options.mode.chained_assignment = None # default='warn'

train_df['day_of_week'] = train_df.date.dt.weekday_name

train_df.tail()
```

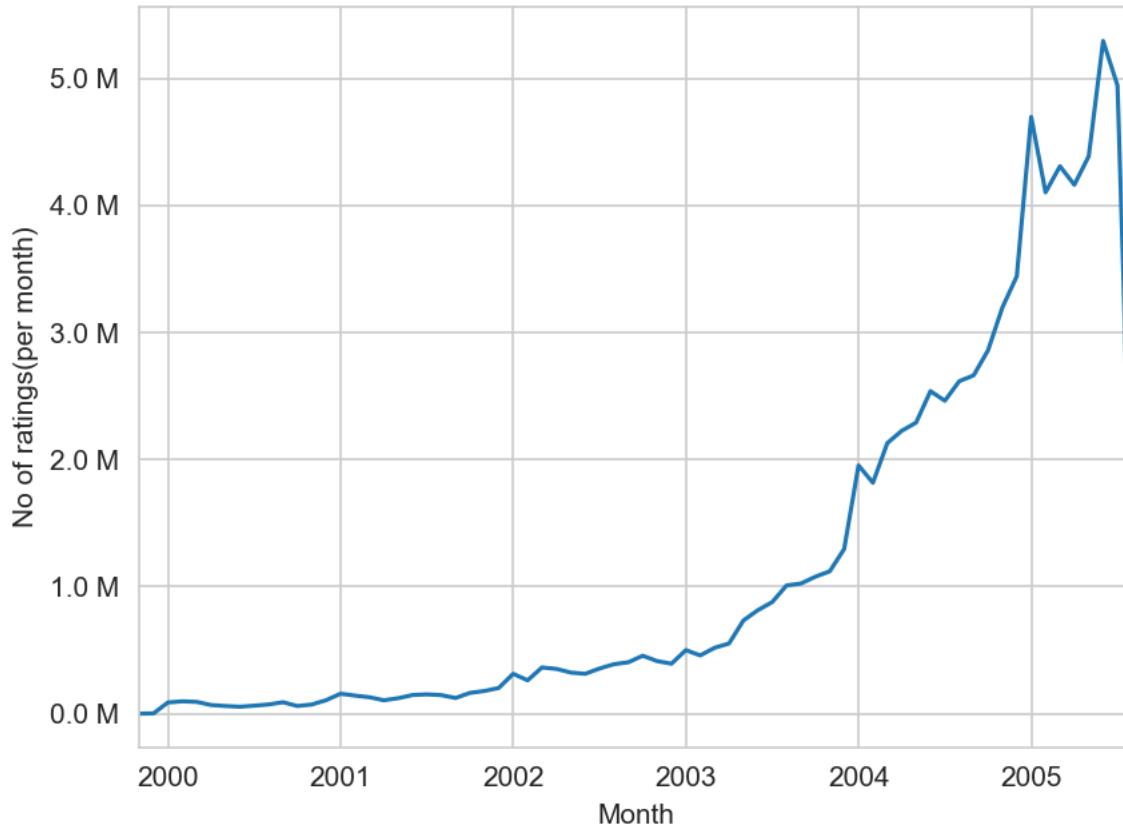
|                 | movie | user    | rating | date       | day_of_week |
|-----------------|-------|---------|--------|------------|-------------|
| <b>80384400</b> | 12074 | 2033618 | 4      | 2005-08-08 | Monday      |
| <b>80384401</b> | 862   | 1797061 | 3      | 2005-08-08 | Monday      |
| <b>80384402</b> | 10986 | 1498715 | 5      | 2005-08-08 | Monday      |
| <b>80384403</b> | 14861 | 500016  | 4      | 2005-08-08 | Monday      |
| <b>80384404</b> | 5926  | 1044015 | 5      | 2005-08-08 | Monday      |

### 3.3.2 Number of Ratings per a month

```
ax = train_df.resample('m', on='date')['rating'].count().plot()
ax.set_title('No of ratings per month (Training data)')
plt.xlabel('Month')
plt.ylabel('No of ratings(per month)')
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
plt.show()
```



No of ratings per month (Training data)



Double-click (or enter) to edit

### 3.3.3 Analysis on the Ratings given by user

```
no_of_rated_movies_per_user = train_df.groupby(by='user')['rating'].count().sort_values(ascen
no_of_rated_movies_per_user.head()
```



user

```
fig = plt.figure(figsize=plt.figaspect(.5))

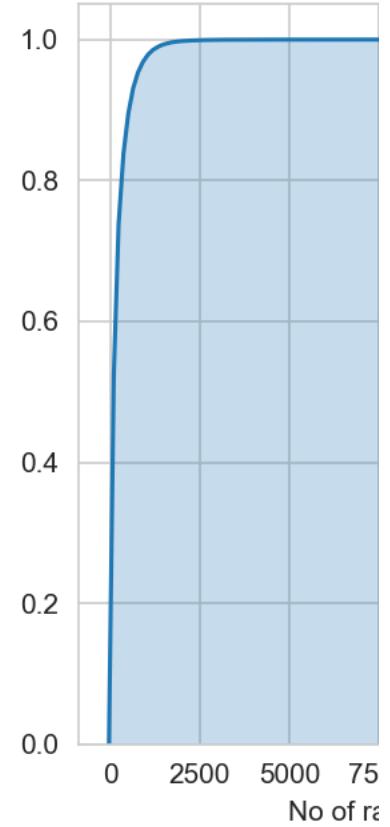
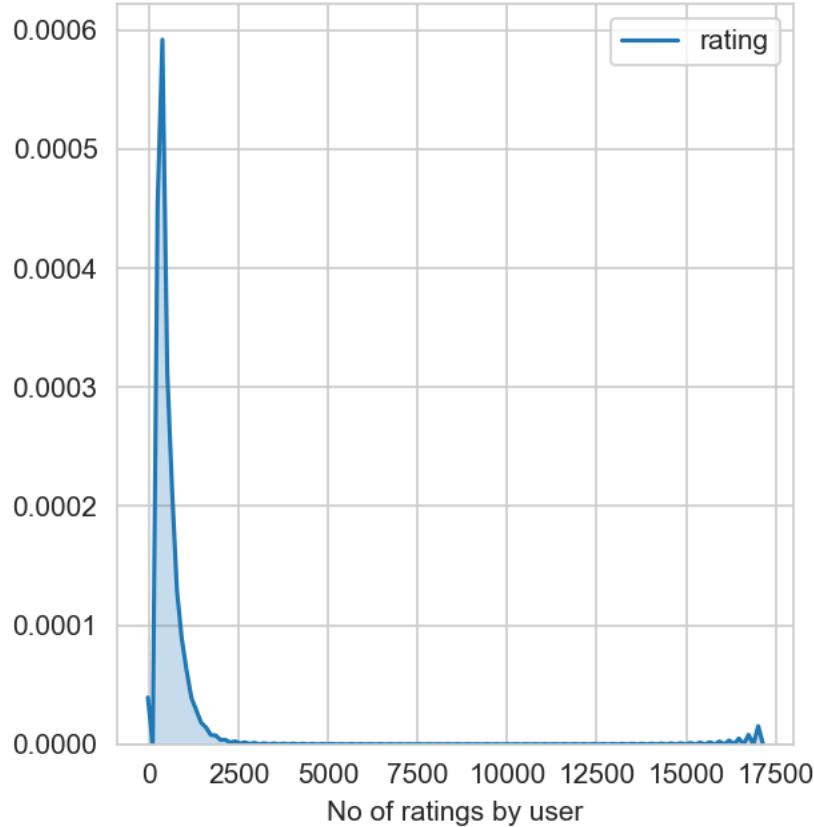
ax1 = plt.subplot(121)
sns.kdeplot(no_of_rated_movies_per_user, shade=True, ax=ax1)
plt.xlabel('No of ratings by user')
plt.title("PDF")

ax2 = plt.subplot(122)
sns.kdeplot(no_of_rated_movies_per_user, shade=True, cumulative=True, ax=ax2)
plt.xlabel('No of ratings by user')
plt.title('CDF')

plt.show()
```



PDF



```
no_of_rated_movies_per_user.describe()
```



```
count    405041.000000
mean      198.459921
```

*There, is something interesting going on with the quantiles..*

```
50%      89.000000
quantiles = no_of_rated_movies_per_user.quantile(np.arange(0,1.01,0.01), interpolation='high
Name: rating, dtype: float64
plt.title("Quantiles and their Values")
quantiles.plot()
# quantiles with 0.05 difference
plt.scatter(x=quantiles.index[::5], y=quantiles.values[::5], c='orange', label="quantiles with
# quantiles with 0.25 difference
plt.scatter(x=quantiles.index[::25], y=quantiles.values[::25], c='m', label = "quantiles with
plt.ylabel('No of ratings by user')
plt.xlabel('Value at the quantile')
plt.legend(loc='best')

# annotate the 25th, 50th, 75th and 100th percentile values....
for x,y in zip(quantiles.index[::25], quantiles[::25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500)
                 ,fontweight='bold')

plt.show()
```



```
quantiles[::5]
```

0.00 1  
0.05 7  
0.10 15  
0.15 21  
0.20 27  
0.25 34  
0.30 41  
0.35 50  
0.40 60  
0.45 73  
0.50 89  
0.55 109  
0.60 133  
0.65 163  
0.70 199  
0.75 245  
0.80 307  
0.85 392  
0.90 520  
0.95 749  
1.00 17112  
Name: rating, dtype: int64

how many ratings at the last 5% of all ratings??

```
print('\n No of ratings at last 5 percentile : {} \n'.format(sum(no_of_rated_movies_per_user>=
```

No of ratings at last 5 percentile : 20305

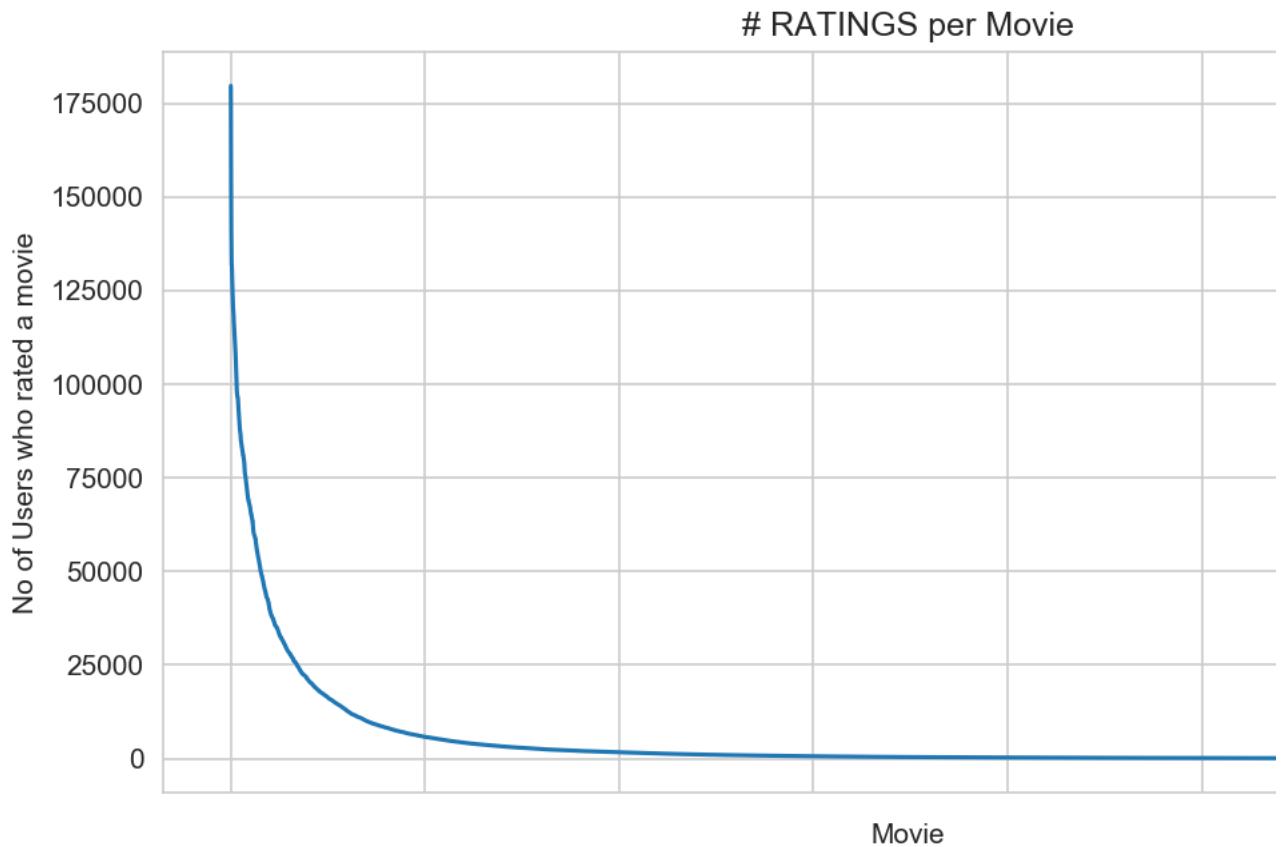
### 3.3.4 Analysis of ratings of a movie given by a user

```
no_of_ratings_per_movie = train_df.groupby(by='movie')['rating'].count().sort_values(ascending=False)

fig = plt.figure(figsize=plt.figaspect(.5))
ax = plt.gca()
plt.plot(no_of_ratings_per_movie.values)
plt.title('# RATINGS per Movie')
plt.xlabel('Movie')
plt.ylabel('No of Users who rated a movie')
ax.set_xticklabels([])

plt.show()
```





- **It is very skewed.. just like number of ratings given per user.**

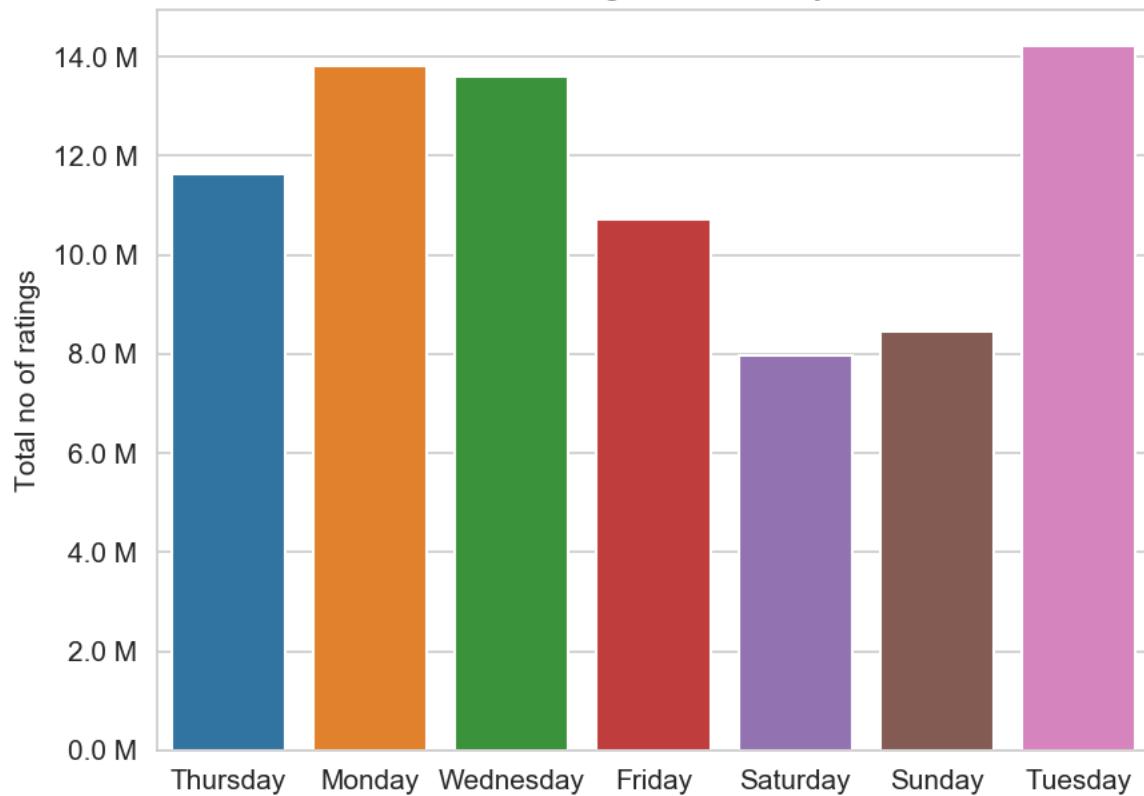
- There are some movies (which are very popular) which are rated by huge number of users.
- But most of the movies (like 90%) got some hundreds of ratings.

### 3.3.5 Number of ratings on each day of the week

```
fig, ax = plt.subplots()
sns.countplot(x='day_of_week', data=train_df, ax=ax)
plt.title('No of ratings on each day...')
plt.ylabel('Total no of ratings')
plt.xlabel('')
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
plt.show()
```

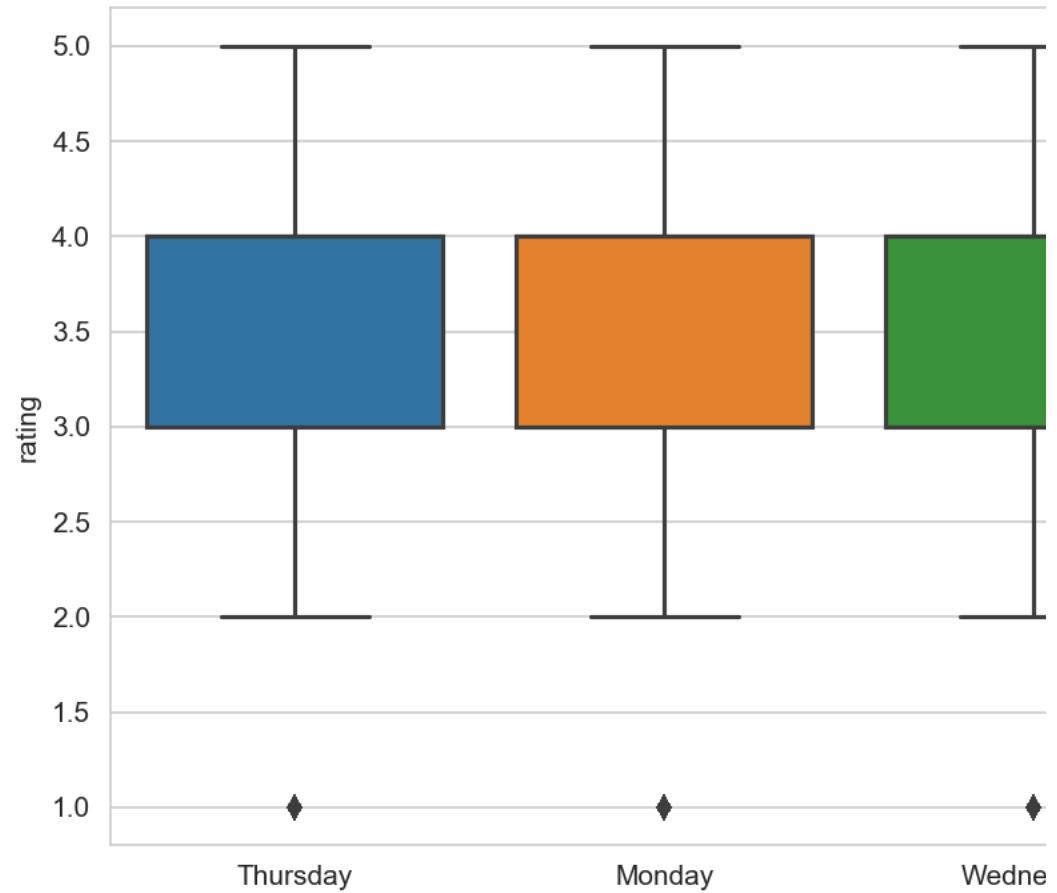


No of ratings on each day...



```
start = datetime.now()
fig = plt.figure(figsize=plt.figaspect(.45))
sns.boxplot(y='rating', x='day_of_week', data=train_df)
plt.show()
print(datetime.now() - start)
```





0:00:25.516197

```
avg_week_df = train_df.groupby(by=[ 'day_of_week' ])[ 'rating' ].mean()
print(" Average ratings")
print("-"*30)
print(avg_week_df)
print("\n")
```



AVerage ratings

Double-click (or enter) to edit

### 3.3.6 Creating sparse matrix from data frame

thurday

0.002403



Name: rating, dtype: float64

#### 3.3.6.1 Creating sparse matrix from train data frame

```
start = datetime.now()
if os.path.isfile('train_sparse_matrix.npz'):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    train_sparse_matrix = sparse.load_npz('train_sparse_matrix.npz')
    print("DONE..")
else:
    print("We are creating sparse_matrix from the dataframe..")
    # create sparse_matrix and store it for after usage.
    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
    # It should be in such a way that, MATRIX[row, col] = data
    train_sparse_matrix = sparse.csr_matrix((train_df.rating.values, (train_df.user.values,
                                                                     train_df.movie.values)),)

    print('Done. It\'s shape is : (user, movie) : ',train_sparse_matrix.shape)
    print('Saving it into disk for furthur usage..')
    # save it into disk
    sparse.save_npz("train_sparse_matrix.npz", train_sparse_matrix)
    print('Done..\n')

print(datetime.now() - start)
```



We are creating sparse\_matrix from the dataframe..  
 Done. It's shape is : (user, movie) : (2649430, 17771)  
 Saving it into disk for furthur usage..  
 Done..

0:02:05.678544

### The Sparsity of Train Sparse Matrix

```
us,mv = train_sparse_matrix.shape
elem = train_sparse_matrix.count_nonzero()

print("Sparsity Of Train matrix : {} % ".format( (1-(elem/(us*mv))) * 100) )

User Sparsity Of Train matrix : 99.8292709259195 %
```

### 3.3.6.2 Creating sparse matrix from test data frame

```

start = datetime.now()
if os.path.isfile('test_sparse_matrix.npz'):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    test_sparse_matrix = sparse.load_npz('test_sparse_matrix.npz')
    print("DONE..")
else:
    print("We are creating sparse_matrix from the dataframe..")
    # create sparse_matrix and store it for after usage.
    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
    # It should be in such a way that, MATRIX[row, col] = data
    test_sparse_matrix = sparse.csr_matrix((test_df.rating.values, (test_df.user.values,
                                                                    test_df.movie.values)))

    print('Done. It\'s shape is : (user, movie) : ',test_sparse_matrix.shape)
    print('Saving it into disk for furthur usage..')
    # save it into disk
    sparse.save_npz("test_sparse_matrix.npz", test_sparse_matrix)
    print('Done..\n')

print(datetime.now() - start)

```

 We are creating sparse\_matrix from the dataframe..  
 Done. It's shape is : (user, movie) : (2649430, 17771)  
 Saving it into disk for furthur usage..  
 Done..

0:00:32.513348

### The Sparsity of Test data Matrix

```

us,mv = test_sparse_matrix.shape
elem = test_sparse_matrix.count_nonzero()

print("Sparsity Of Test matrix : {} % ".format( (1-(elem/(us*mv))) * 100) )

```

 Sparsity Of Test matrix : 99.95731772988694 %

### 3.3.7 Finding Global average of all movie ratings, Average rating per user, and Average rating per movie

```

# get the user averages in dictionary (key: user_id/movie_id, value: avg rating)

def get_average_ratings(sparse_matrix, of_users):

```

# average ratings of user/movies

```
# average ratings of user/axes
ax = 1 if of_users else 0 # 1 - User axes, 0 - Movie axes

# ".A1" is for converting Column_Matrix to 1-D numpy array
sum_of_ratings = sparse_matrix.sum(axis=ax).A1
# Boolean matrix of ratings ( whether a user rated that movie or not)
is_rated = sparse_matrix!=0
# no of ratings that each user OR movie..
no_of_ratings = is_rated.sum(axis=ax).A1

# max_user and max_movie ids in sparse matrix
u,m = sparse_matrix.shape
# create a dictionary of users and their average ratings..
average_ratings = { i : sum_of_ratings[i]/no_of_ratings[i]
                     for i in range(u if of_users else m)
                     if no_of_ratings[i] !=0}

# return that dictionary of average ratings
return average_ratings
```

### 3.3.7.1 finding global average of all movie ratings

```
train_averages = dict()
# get the global average of ratings in our train set.
train_global_average = train_sparse_matrix.sum()/train_sparse_matrix.count_nonzero()
train_averages['global'] = train_global_average
train_averages
```

 {'global': 3.582890686321557}

### 3.3.7.2 finding average rating per user

```
train_averages['user'] = get_average_ratings(train_sparse_matrix, of_users=True)
print('\nAverage rating of user 10 :',train_averages['user'][10])
```

 Average rating of user 10 : 3.3781094527363185

### 3.3.7.3 finding average rating per movie

```
train_averages['movie'] = get_average_ratings(train_sparse_matrix, of_users=False)
print('\n Average rating of movie 15 :',train_averages['movie'][15])
```

 Average rating of movie 15 : 3.3038461538461537

Double-click (or enter) to edit

### 3.3.7.4 PDF's & CDF's of Avg.Ratings of Users & Movies (In Train Data)

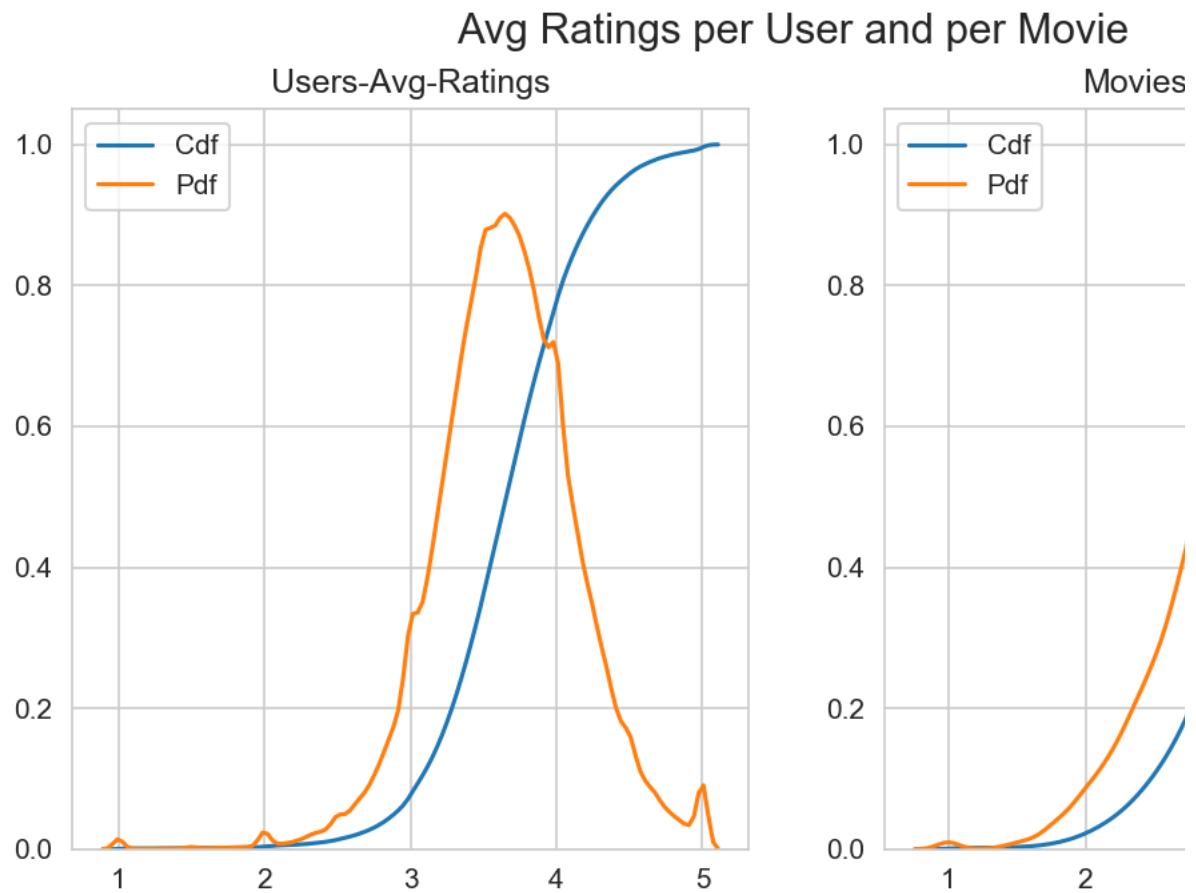
```
start = datetime.now()
# draw pdfs for average rating per user and average
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))
fig.suptitle('Avg Ratings per User and per Movie', fontsize=15)

ax1.set_title('Users-Avg-Ratings')
# get the list of average user ratings from the averages dictionary..
user_averages = [rat for rat in train_averages['user'].values()]
sns.distplot(user_averages, ax=ax1, hist=False,
             kde_kws=dict(cumulative=True), label='Cdf')
sns.distplot(user_averages, ax=ax1, hist=False, label='Pdf')

ax2.set_title('Movies-Avg-Rating')
# get the list of movie_average_ratings from the dictionary..
movie_averages = [rat for rat in train_averages['movie'].values()]
sns.distplot(movie_averages, ax=ax2, hist=False,
             kde_kws=dict(cumulative=True), label='Cdf')
sns.distplot(movie_averages, ax=ax2, hist=False, label='Pdf')

plt.show()
print(datetime.now() - start)
```





0:02:36.494092

### 3.3.8 Cold Start problem

#### 3.3.8.1 Cold Start problem with Users

```
total_users = len(np.unique(df.user))
users_train = len(train_averages['user'])
new_users = total_users - users_train

print('\nTotal number of Users  :', total_users)
print('\nNumber of Users in Train data :', users_train)
print("\nNo of Users that didn't appear in train data: {}({} %) \n ".format(new_users,
                                                                           np.round((new_users/t
```



We might have to handle **new users ( 75148 )** who didn't appear in train data.

### 3.3.8.2 Cold Start problem with Movies

```
total_movies = len(np.unique(df.movie))
movies_train = len(train_averages['movie'])
new_movies = total_movies - movies_train

print('\nTotal number of Movies :', total_movies)
print('\nNumber of Users in Train data :', movies_train)
print("\nNo of Movies that didn't appear in train data: {}({} %) \n ".format(new_movies,
                                                                           np.round((new_movies/
```



Total number of Movies : 17770

Number of Users in Train data : 17424

No of Movies that didn't appear in train data: 346(1.95 %)

We might have to handle **346 movies** (small comparatively) in test data

Double-click (or enter) to edit

## 3.4 Computing Similarity matrices

### 3.4.1 Computing User-User Similarity matrix

1. Calculating User User Similarity\_Matrix is **not very easy**(unless you have huge Computing Power users being late).

- You can try if you want to. Your system could crash or the program stops with **Memory Error**

#### 3.4.1.1 Trying with all dimensions (17k dimensions per user)

```
from sklearn.metrics.pairwise import cosine_similarity
```

```
def compute_user_similarity(sparse_matrix, compute_for_few=False, top = 100, verbose=False, v
```

```

        draw_time_taken=True):
no_of_users, _ = sparse_matrix.shape
# get the indices of non zero rows(users) from our sparse matrix
row_ind, col_ind = sparse_matrix.nonzero()
row_ind = sorted(set(row_ind)) # we don't have to
time_taken = list() # time taken for finding similar users for an user..

# we create rows, cols, and data lists.., which can be used to create sparse matrices
rows, cols, data = list(), list(), list()
if verbose: print("Computing top",top,"similarities for each user..")

start = datetime.now()
temp = 0

for row in row_ind[:top] if compute_for_few else row_ind:
    temp = temp+1
    prev = datetime.now()

    # get the similarity row for this user with all other users
    sim = cosine_similarity(sparse_matrix.getrow(row), sparse_matrix).ravel()
    # We will get only the top ''top'' most similar users and ignore rest of them..
    top_sim_ind = sim.argsort()[-top:]
    top_sim_val = sim[top_sim_ind]

    # add them to our rows, cols and data
    rows.extend([row]*top)
    cols.extend(top_sim_ind)
    data.extend(top_sim_val)
    time_taken.append(datetime.now().timestamp() - prev.timestamp())
    if verbose:
        if temp%verb_for_n_rows == 0:
            print("computing done for {} users [ time elapsed : {} ]"
                  .format(temp, datetime.now()-start))

# lets create sparse matrix out of these and return it
if verbose: print('Creating Sparse matrix from the computed similarities')
#return rows, cols, data

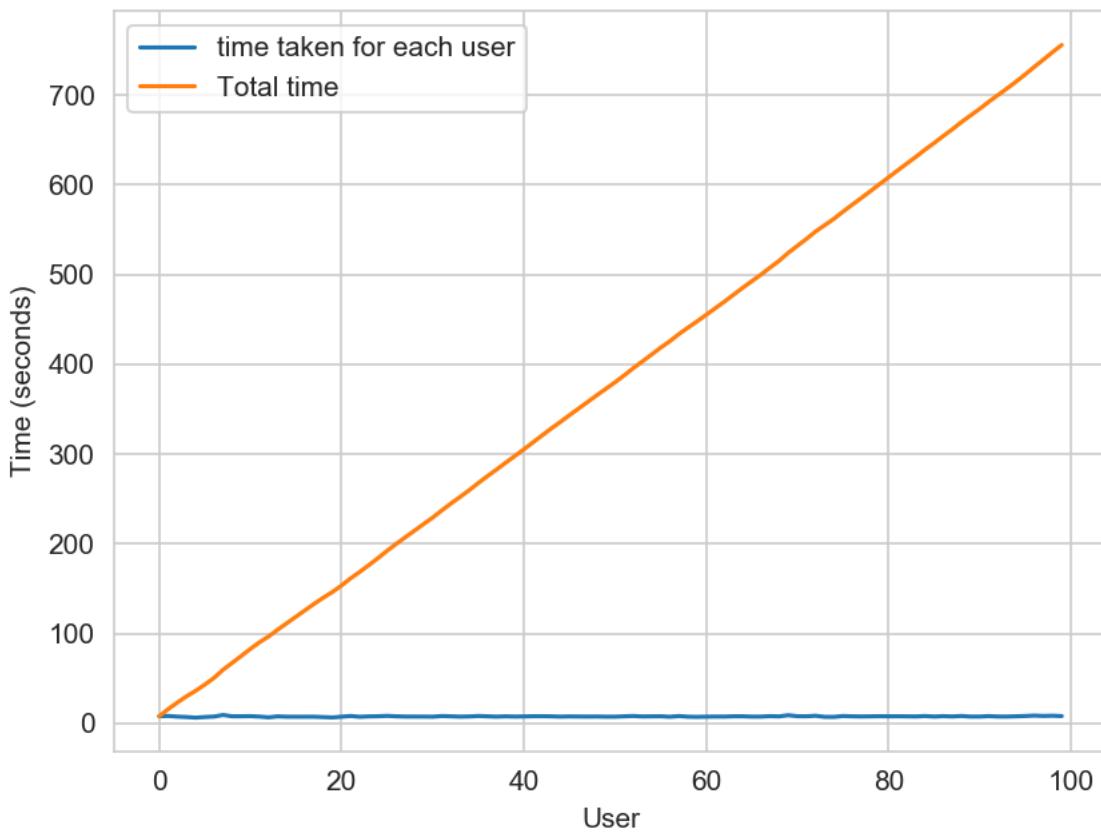
if draw_time_taken:
    plt.plot(time_taken, label = 'time taken for each user')
    plt.plot(np.cumsum(time_taken), label='Total time')
    plt.legend(loc='best')
    plt.xlabel('User')
    plt.ylabel('Time (seconds)')
    plt.show()

return sparse.csr_matrix((data, (rows, cols)), shape=(no_of_users, no_of_users)), time_ta

```

```
print("-"*100)
print("Time taken :",datetime.now()-start)
```

Computing top 100 similarities for each user..  
 computing done for 20 users [ time elapsed : 0:02:26.186259 ]  
 computing done for 40 users [ time elapsed : 0:04:57.600051 ]  
 computing done for 60 users [ time elapsed : 0:07:27.747405 ]  
 computing done for 80 users [ time elapsed : 0:10:00.443105 ]  
 computing done for 100 users [ time elapsed : 0:12:35.667537 ]  
 Creating Sparse matrix from the computed similarities




---

Time taken : 0:12:58.731790

### 3.4.1.2 Trying with reduced dimensions (Using TruncatedSVD for dimensionality reduction)

Double-click (or enter) to edit

- We have **405,041 users** in our training set and computing similarities between them..( **17K dimensions** )
- From above plot, It took roughly **8.88 sec** for computing similar users for **one user**
- We have **405,041 users** with us in training set.
- $405041 \times 8.88 = 3596764.08 \text{ sec} = 59946.068 \text{ min} = 999.101133333 \text{ hours} = 41.62 \text{ days}$

- Even if we run on 4 cores parallelly (a typical system now a days), It will still take almost 1000 days

IDEA: Instead, we will try to reduce the dimensions using SVD, so that **it might** speed up the process

```
from datetime import datetime
from sklearn.decomposition import TruncatedSVD

start = datetime.now()

# initialize the algorithm with some parameters..
# All of them are default except n_components. n_iter is for Randomized SVD solver.
netflix_svd = TruncatedSVD(n_components=500, algorithm='randomized', random_state=15)
trunc_svd = netflix_svd.fit_transform(train_sparse_matrix)

print(datetime.now()-start)
```

0:57:33.677420

Here,

- $\sum$   $\leftarrow$  (netflix\_svd.singular\_values\_)
- $V^T$   $\leftarrow$  (netflix\_svd.components\_)
- $U$  is not returned. instead **Projection\_of\_X** onto the new vectorspace is returned.
- It uses **randomized svd** internally, which returns **All 3 of them separately**. Use that instead..

```
expl_var = np.cumsum(netflix_svd.explained_variance_ratio_)

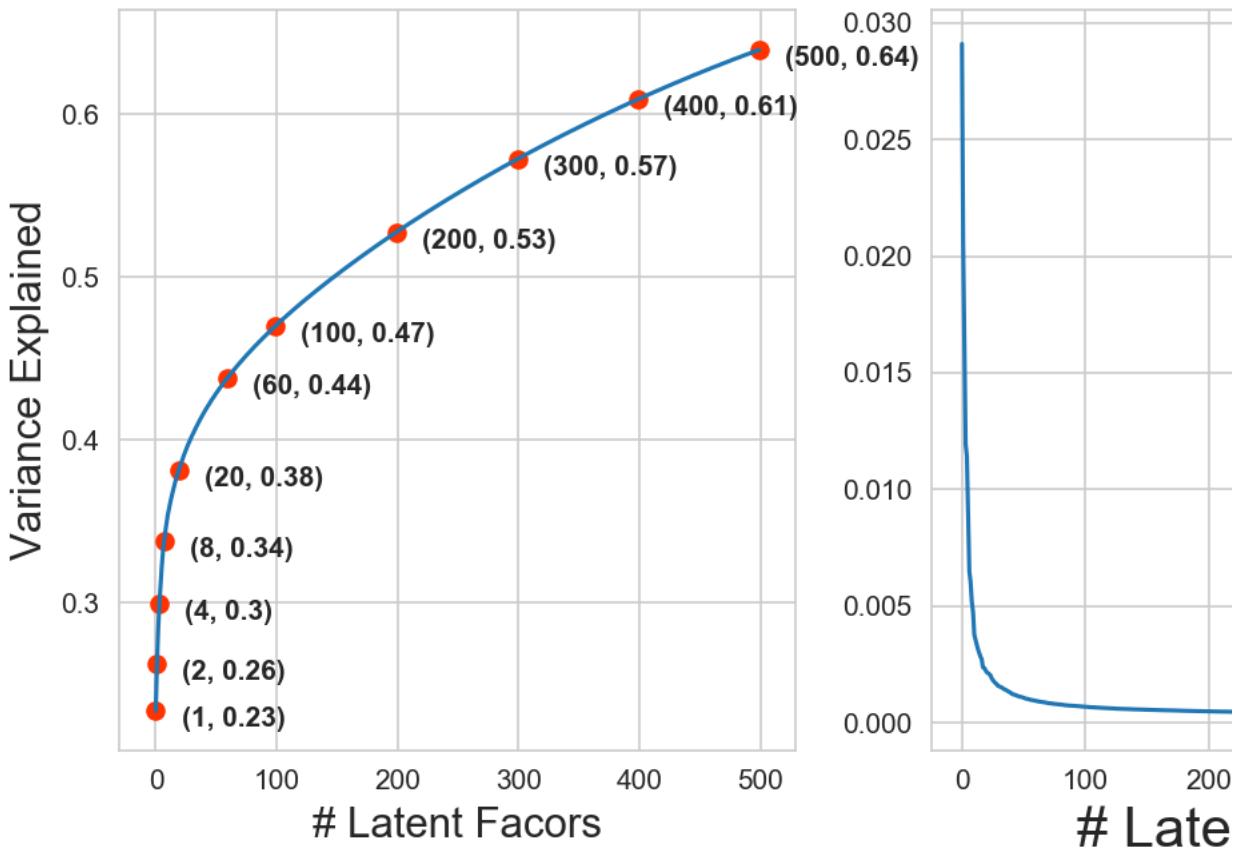
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))

ax1.set_ylabel("Variance Explained", fontsize=15)
ax1.set_xlabel("# Latent Factors", fontsize=15)
ax1.plot(expl_var)
# annotate some (latentfactors, expl_var) to make it clear
ind = [1, 2, 4, 8, 20, 60, 100, 200, 300, 400, 500]
ax1.scatter(x = [i-1 for i in ind], y = expl_var[[i-1 for i in ind]], c='ff3300')
for i in ind:
    ax1.annotate(s = "({}, {})".format(i, np.round(expl_var[i-1], 2)), xy=(i-1, expl_var[i-1])
                 xytext = (i+20, expl_var[i-1] - 0.01), fontweight='bold')

change_in_expl_var = [expl_var[i+1] - expl_var[i] for i in range(len(expl_var)-1)]
ax2.plot(change_in_expl_var)

ax2.set_ylabel("Gain in Var_Expl with One Additional LF", fontsize=10)
ax2.yaxis.set_label_position("right")
ax2.set_xlabel("# Latent Factors", fontsize=20)
```

```
plt.show()
```



```
for i in ind:
    print("({}, {})".format(i, np.round(expl_var[i-1], 2)))
```



```
(1, 0.23)
(2, 0.26)
(4, 0.3)
(8, 0.34)
(20, 0.38)
(60, 0.44)
(100, 0.47)
(200, 0.53)
(300, 0.57)
(400, 0.61)
(500, 0.64)
```

I think 500 dimensions is good enough

- By just taking **(20 to 30)** latent factors, explained variance that we could get is **20 %**.
- To take it to **60%**, we have to take **almost 400 latent factors**. It is not fare.
- It basically is the **gain of variance explained**, if we **add one additional latent factor to it**.

- By adding one by one latent factor to it, the **gain in explained variance** with that addition is sorted that way).
- **LHS Graph:**
  - **x** --- ( No of latent factors ),
  - **y** --- ( The variance explained by taking x latent factors)
- **More decrease in the line (RHS graph) :**
  - We are getting more explained variance than before.
- **Less decrease in that line (RHS graph) :**
  - We are not getting benefitted from adding latent factor further. This is what is shown in the
- **RHS Graph:**
  - **x** --- ( No of latent factors ),
  - **y** --- ( Gain in Expl\_Var by taking one additional latent factor)

```
# Let's project our Original U_M matrix into 500 Dimensional space...
start = datetime.now()
trunc_matrix = train_sparse_matrix.dot(netflix_svd.components_.T)
print(datetime.now() - start)
```

 0:00:47.090264

```
type(trunc_matrix), trunc_matrix.shape
```

 (numpy.ndarray, (2649430, 500))

- Let's convert this to actual sparse matrix and store it for future purposes

```
if not os.path.isfile('trunc_sparse_matrix.npz'):
    # create that sparse sparse matrix
    trunc_sparse_matrix = sparse.csr_matrix(trunc_matrix)
    # Save this truncated sparse matrix for later usage..
    sparse.save_npz('trunc_sparse_matrix', trunc_sparse_matrix)
else:
    trunc_sparse_matrix = sparse.load_npz('trunc_sparse_matrix.npz')
```

```
trunc_sparse_matrix.shape
```

 (2649430, 500)

```
start = datetime.now()
trunc_u_u_sim_matrix, _ = compute_user_similarity(trunc_sparse_matrix, compute_for_few=True,
                                                 verb_for_n_rows=10)
```

```

print("-"*50)
print("time:",datetime.now()-start)

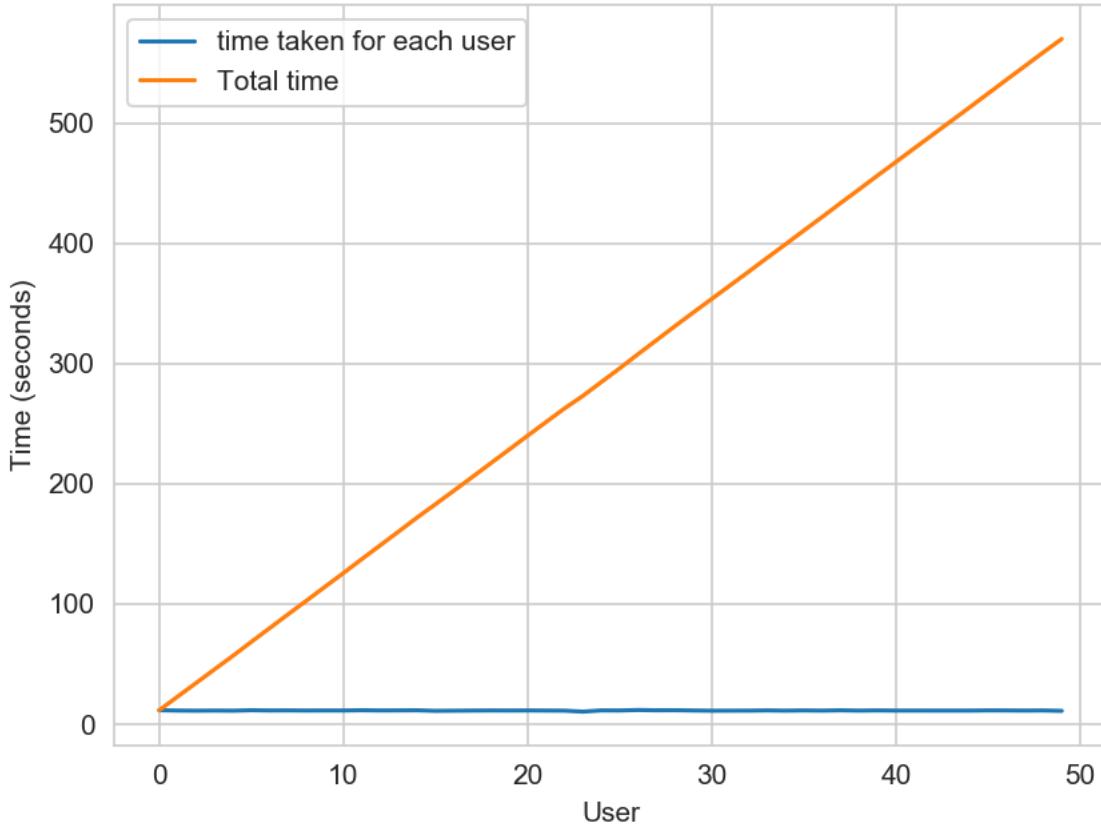


- ➊ Computing top 50 similarities for each user..
    computing done for 10 users [ time elapsed : 0:01:54.375445 ]
    computing done for 20 users [ time elapsed : 0:03:48.634715 ]
    computing done for 30 users [ time elapsed : 0:05:42.733062 ]
    computing done for 40 users [ time elapsed : 0:07:36.623472 ]
    computing done for 50 users [ time elapsed : 0:09:30.260668 ]



Creating Sparse matrix from the computed similarities


```



-----  
time: 0:10:28.489709

: This is taking more time for each user than Original one.

- from above plot, It took almost **12.18** for computing similar users for **one user**
- We have **405041 users** with us in training set.
- $405041 \times 12.18 = 4933399.38 \text{ sec} = 82223.323 \text{ min} = 1370.38871 \text{ h} = 57.099529861 \text{ days...}$ 
  - Even we run on 4 cores parallelly (a typical system now a days), It will still take almost ...
- **Why did this happen...??**

- Just think about it. It's not that difficult.

----- ( sparse & dense.....get it ?? ) -----

## Is there any other way to compute user user similarity..??

-An alternative is to compute similar users for a particular user, whenever required (**ie., Run time**) - Which tells us whether we already computed or not.. - **If not** : - Compute top (let's just say, 1000) most similar users to our datastructure, so that we can just access it(similar users) without recomputing it again. - - directly from our datastructure, which has that information. - In production time, We might have to recalculate this information long time ago. Because user preferences changes over time. If we could maintain some kind of Time (recompute it). - - **Which datastructure to use**: - It is purely implementation dependant. - One simple approach is to use Dictionaries. - - **key** : *userid* - **value**: *Again a dictionary* - **key** : *Similar User* - **value**: *Similarity Value*

### 3.4.2 Computing Movie-Movie Similarity matrix

```
start = datetime.now()
if not os.path.isfile('m_m_sim_sparse.npz'):
    print("It seems you don't have that file. Computing movie_movie similarity...")
    start = datetime.now()
    m_m_sim_sparse = cosine_similarity(X=train_sparse_matrix.T, dense_output=False)
    print("Done..")
    # store this sparse matrix in disk before using it. For future purposes.
    print("Saving it to disk without the need of re-computing it again.. ")
    sparse.save_npz("m_m_sim_sparse.npz", m_m_sim_sparse)
    print("Done..")
else:
    print("It is there, We will get it.")
    m_m_sim_sparse = sparse.load_npz("m_m_sim_sparse.npz")
    print("Done ...")

print("It's a ",m_m_sim_sparse.shape," dimensional matrix")

print(datetime.now() - start)
```

 It seems you don't have that file. Computing movie\_movie similarity...
 Done..
 Saving it to disk without the need of re-computing it again..
 Done..
 It's a (17771, 17771) dimensional matrix
 0:15:21.738551

`m_m_sim_sparse.shape`



(17771, 17771)

- Even though we have similarity measure of each movie, with all other movies, We generally don't
- Most of the times, only top\_xxx similar items matters. It may be 10 or 100.
- We take only those top similar movie ratings and store them in a saperate dictionary.

```
movie_ids = np.unique(m_m_sim_sparse.nonzero()[1])
```

```
start = datetime.now()
similar_movies = dict()
for movie in movie_ids:
    # get the top similar movies and store them in the dictionary
    sim_movies = m_m_sim_sparse[movie].toarray().ravel().argsort()[:-1][1:]
    similar_movies[movie] = sim_movies[:100]
print(datetime.now() - start)

# just testing similar movies for movie_15
similar_movies[15]
```



0:00:51.829120

```
array([ 8279,  8013, 16528,  5927, 13105, 12049,  4424, 10193, 17590,
       4549,  3755,  590, 14059, 15144, 15054,  9584,  9071,  6349,
      16402,  3973,  1720,  5370, 16309,  9376,  6116,  4706,  2818,
      778, 15331, 1416, 12979, 17139, 17710,  5452,  2534,   164,
     15188,  8323, 2450, 16331,  9566, 15301, 13213, 14308, 15984,
     10597,  6426,  5500,  7068,  7328,  5720,  9802,   376, 13013,
      8003, 10199, 3338, 15390,  9688, 16455, 11730,  4513,   598,
     12762, 2187,   509,  5865,  9166, 17115, 16334,  1942,  7282,
     17584, 4376, 8988, 8873,  5921,  2716, 14679, 11947, 11981,
      4649,   565, 12954, 10788, 10220, 10963,  9427,  1690,  5107,
      7859,  5969, 1510,  2429,   847,  7845,  6410, 13931,  9840,
     3706], dtype=int64)
```

Double-click (or enter) to edit

### 3.4.3 Finding most similar movies using similarity matrix

Does Similarity really works as the way we expected...?\_\_

Let's pick some random movie and check for its similar movies....

```
# First Let's load the movie details into soe datafame..
# movie details are in 'netflix/movie_titles.csv'
```

```
movie_titles = pd.read_csv("movie_titles.csv", sep=',', header = None,
```

```
names=['movie_id', 'year_of_release', 'title'], verbose=True,
index_col = 'movie_id', encoding = "ISO-8859-1")
```

```
movie_titles.head()
```

Tokenization took: 15.09 ms  
 Type conversion took: 29.15 ms  
 Parser memory cleanup took: 0.00 ms

| movie_id | year_of_release | title                        |
|----------|-----------------|------------------------------|
| 1        | 2003.0          | Dinosaur Planet              |
| 2        | 2004.0          | Isle of Man TT 2004 Review   |
| 3        | 1997.0          | Character                    |
| 4        | 1994.0          | Paula Abdul's Get Up & Dance |
| 5        | 2004.0          | The Rise and Fall of ECW     |

## Similar Movies for 'Vampire Journals'

```
mv_id = 67
```

```
print("\nMovie ---->, movie_titles.loc[mv_id].values[1])
```

```
print("\nIt has {} Ratings from users.".format(train_sparse_matrix[:,mv_id].getnnz()))
```

```
print("\nWe have {} movies which are similar to this and we will get only top most..".format(
```

Movie ----> Vampire Journals

It has 270 Ratings from users.

We have 17284 movies which are similar to this and we will get only top most..

```
similarities = m_m_sim_sparse[mv_id].toarray().ravel()
```

```
similar_indices = similarities.argsort()[:-1][1:]
```

```
similarities[similar_indices]
```

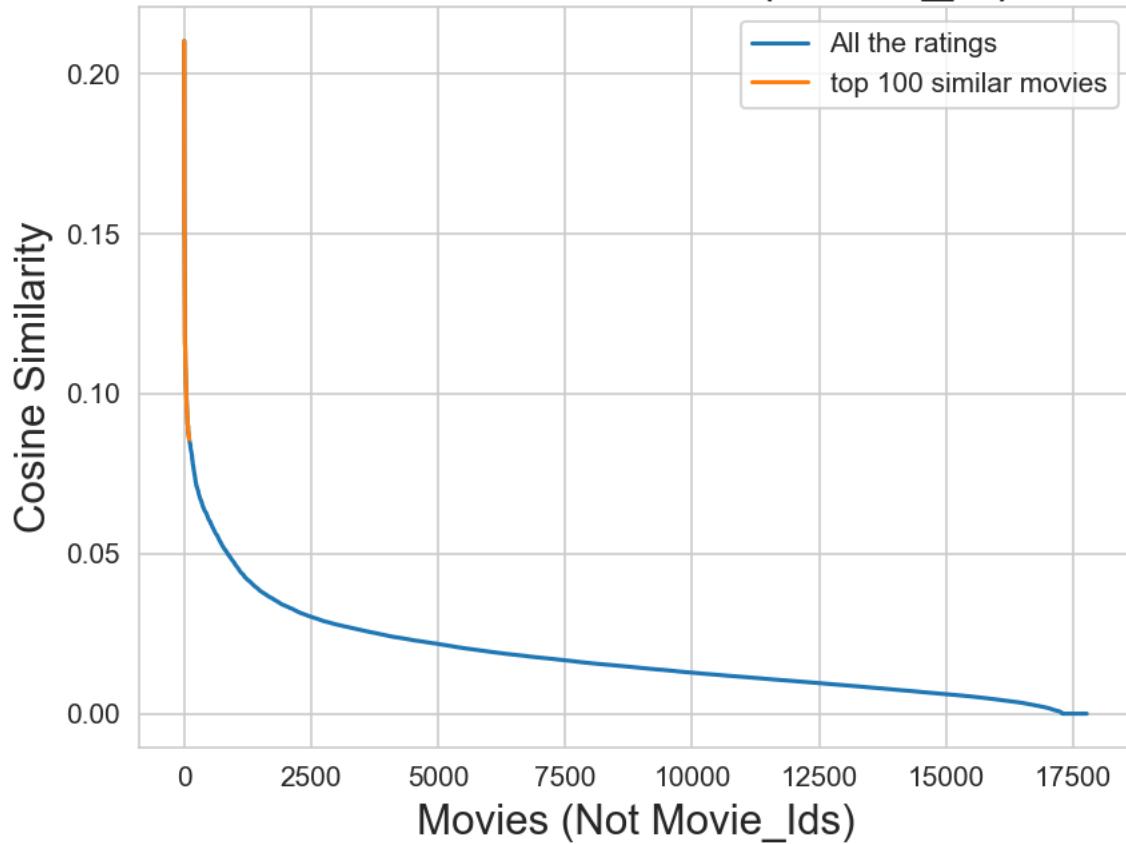
```
sim_indices = similarities.argsort()[:-1][1:] # It will sort and reverse the array and ignore
# and return its indices(movie_ids)
```

```
plt.plot(similarities[sim_indices], label='All the ratings')
plt.plot(similarities[sim_indices[:100]], label='top 100 similar movies')
plt.title("Similar Movies of {}(movie_id)".format(mv_id), fontsize=20)
plt.xlabel("Movies (Not Movie_Ids)", fontsize=15)
```

```
plt.ylabel("Cosine Similarity", fontsize=15)
plt.legend()
plt.show()
```



## Similar Movies of 67(movie\_id)



Double-click (or enter) to edit

Double-click (or enter) to edit

### Top 10 similar movies

```
movie_titles.loc[sim_indices[:10]]
```



| movie_id | year_of_release | title                    |
|----------|-----------------|--------------------------|
| 323      | 1999.0          | Modern Vampires          |
| 4044     | 1998.0          | Subspecies 4: Bloodstorm |
| 1688     | 1993.0          | To Sleep With a Vampire  |
| 13962    | 2001.0          | Dracula: The Dark Prince |
| 12053    | 1993.0          | Dracula Rising           |
| 16279    | 2002.0          | Vampires: Los Muertos    |
| 4667     | 1996.0          | Vampirella               |
| 1900     | 1997.0          | Club Vampire             |
| 13873    | 2001.0          | The Breed                |
| 15867    | 2003.0          | Dracula II: Ascension    |

Double-click (or enter) to edit

Similarly, we can **find similar users** and compare how similar they are.

Double-click (or enter) to edit

Double-click (or enter) to edit

## 4. Machine Learning Models



```
def get_sample_sparse_matrix(sparse_matrix, no_users, no_movies, path, verbose = True):
    """
        It will get it from the ''path'' if it is present or It will create
        and store the sampled sparse matrix in the path specified.
    """

    # get (row, col) and (rating) tuple from sparse_matrix...
    row_ind, col_ind, ratings = sparse.find(sparse_matrix)
    users = np.unique(row_ind)
    movies = np.unique(col_ind)
```

```

print("Original Matrix : (users, movies) -- ({} {})".format(len(users), len(movies)))
print("Original Matrix : Ratings -- {}".format(len(ratings)))

# It just to make sure to get same sample everytime we run this program..
# and pick without replacement....
np.random.seed(15)
sample_users = np.random.choice(users, no_users, replace=False)
sample_movies = np.random.choice(movies, no_movies, replace=False)
# get the boolean mask or these sampled_items in originl row/col_inds..
mask = np.logical_and( np.isin(row_ind, sample_users),
                       np.isin(col_ind, sample_movies) )

sample_sparse_matrix = sparse.csr_matrix((ratings[mask], (row_ind[mask], col_ind[mask])),
                                           shape=(max(sample_users)+1, max(sample_movies)+1))

if verbose:
    print("Sampled Matrix : (users, movies) -- ({} {})".format(len(sample_users), len(sample_movies)))
    print("Sampled Matrix : Ratings --", format(ratings[mask].shape[0]))

print('Saving it into disk for furthur usage..')
# save it into disk
sparse.save_npz(path, sample_sparse_matrix)
if verbose:
    print('Done..\n')

return sample_sparse_matrix

```

## 4.1 Sampling Data

### 4.1.1 Build sample train data from the train data

```

start = datetime.now()
path = "sample_train_sparse_matrix.npz"
if os.path.isfile(path):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    sample_train_sparse_matrix = sparse.load_npz(path)
    print("DONE..")
else:
    # get 10k users and 1k movies from available data
    sample_train_sparse_matrix = get_sample_sparse_matrix(train_sparse_matrix, no_users=30000
                                                          path = path)

print(datetime.now() - start)

```



```
Original Matrix : (users, movies) -- (405041 17424)
Original Matrix : Ratings -- 80384405

Sampled Matrix : (users, movies) -- (30000 3000)
Sampled Matrix : Ratings -- 1029316
Saving it into disk for furthur usage..
Done..
```

0:01:49.750637

#### 4.1.2 Build sample test data from the test data

```
start = datetime.now()

path = "sample_test_sparse_matrix.npz"
if os.path.isfile(path):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    sample_test_sparse_matrix = sparse.load_npz(path)
    print("DONE..")
else:
    # get 5k users and 500 movies from available data
    sample_test_sparse_matrix = get_sample_sparse_matrix(test_sparse_matrix, no_users=7500, n
                                                       path = "sample_test_sparse_matrix.npz")
print(datetime.now() - start)
```

👤 Original Matrix : (users, movies) -- (349312 17757)
Original Matrix : Ratings -- 20096102

```
Sampled Matrix : (users, movies) -- (7500 750)
Sampled Matrix : Ratings -- 19273
Saving it into disk for furthur usage..
Done..
```

0:00:24.446595

Double-click (or enter) to edit

#### 4.2 Finding Global Average of all movie ratings, Average rating per Use Movie (from sampled train)

```
sample_train_averages = dict()
```

##### 4.2.1 Finding Global Average of all movie ratings

```
# get the global average of ratings in our train set.  
global_average = sample_train_sparse_matrix.sum()/sample_train_sparse_matrix.count_nonzero()  
sample_train_averages['global'] = global_average  
sample_train_averages  
 {'global': 3.5902997718873504}
```

#### 4.2.2 Finding Average rating per User

```
sample_train_averages['user'] = get_average_ratings(sample_train_sparse_matrix, of_users=True)  
print('\nAverage rating of user 1515220 :',sample_train_averages['user'][1515220])
```



Average rating of user 1515220 : 3.923076923076923

#### 4.2.3 Finding Average rating per Movie

```
sample_train_averages['movie'] = get_average_ratings(sample_train_sparse_matrix, of_users=False)  
print('\n Average rating of movie 15153 :',sample_train_averages['movie'][15153])
```



AVerage rating of movie 15153 : 2.7974683544303796

Double-click (or enter) to edit

### 4.3 Featurizing data

```
print('\n No of ratings in Our Sampled train matrix is : {}'.format(sample_train_sparse_matrix.shape[0]))  
print('\n No of ratings in Our Sampled test matrix is : {}'.format(sample_test_sparse_matrix.shape[0]))
```



No of ratings in Our Sampled train matrix is : 1029316

No of ratings in Our Sampled test matrix is : 19273

#### 4.3.1 Featurizing data for regression problem

##### 4.3.1.1 Featurizing train data

```
# get users, movies and ratings from our samples train sparse matrix
```

```

sample_train_users, sample_train_movies, sample_train_ratings = sparse.find(sample_train_spar

start = datetime.now()
if os.path.isfile('reg_train.csv'):
    print("File already exists you don't have to prepare again... ")
else:
    print('preparing {} tuples for the dataset..\n'.format(len(sample_train_ratings)))
    with open('sample/small/reg_train.csv', mode='w') as reg_data_file:
        count = 0
        for (user, movie, rating) in zip(sample_train_users, sample_train_movies, sample_train_
            st = datetime.now()
            # print(user, movie)
            #----- Ratings of "movie" by similar users of "user" -----
            # compute the similar Users of the "user"
            user_sim = cosine_similarity(sample_train_sparse_matrix[user], sample_train_spars
            top_sim_users = user_sim.argsort()[:-1][1:] # we are ignoring 'The User' from it
            # get the ratings of most similar users for this movie
            top_ratings = sample_train_sparse_matrix[top_sim_users, movie].toarray().ravel()
            # we will make it's length "5" by adding movie averages to .
            top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5])
            top_sim_users_ratings.extend([sample_train_averages['movie'][movie]]*(5 - len(top
            # print(top_sim_users_ratings, end=" ")

            #----- Ratings by "user" to similar movies of "movie" -----
            # compute the similar movies of the "movie"
            movie_sim = cosine_similarity(sample_train_sparse_matrix[:,movie].T, sample_train_
            top_sim_movies = movie_sim.argsort()[:-1][1:] # we are ignoring 'The User' from
            # get the ratings of most similar movie rated by this user..
            top_ratings = sample_train_sparse_matrix[user, top_sim_movies].toarray().ravel()
            # we will make it's length "5" by adding user averages to.
            top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:5])
            top_sim_movies_ratings.extend([sample_train_averages['user'][user]]*(5-len(top_si
            # print(top_sim_movies_ratings, end=" : -- ")

#-----prepare the row to be stores in a file-----#
row = list()
row.append(user)
row.append(movie)
# Now add the other features to this data...
row.append(sample_train_averages['global']) # first feature
# next 5 features are similar_users "movie" ratings
row.extend(top_sim_users_ratings)
# next 5 features are "user" ratings for similar_movies
row.extend(top_sim_movies_ratings)
# Avg_user rating
row.append(sample_train_averages['user'][user])
# Avg_movie rating
row.append(sample_train_averages['movie'][movie])

# finalley, The actual Rating of this user-movie pair...
row.append(rating)

```

```

        row.append(rating)
        count = count + 1

        # add rows to the file opened..
        reg_data_file.write(','.join(map(str, row)))
        reg_data_file.write('\n')
        if (count)%10000 == 0:
            # print(','.join(map(str, row)))
            print("Done for {} rows---- {} ".format(count, datetime.now() - start))

print(datetime.now() - start)

```

👤 File already exists you don't have to prepare again...  
0:00:00.000997

## Reading from the file to make a Train\_dataframe

```
reg_train = pd.read_csv('reg_train.csv', names = ['user', 'movie', 'GAvg', 'sur1', 'sur2', 'sur3', 'sur4', 'sur5', 'smr1', 'smr2', 'smr3', 'smr4', 'smr5'])
reg_train.head()
```

|   | user   | movie | GAvg     | sur1 | sur2 | sur3 | sur4 | sur5 | smr1 | smr2 | smr3 | smr4 | smr5 |
|---|--------|-------|----------|------|------|------|------|------|------|------|------|------|------|
| 0 | 53406  | 33    | 3.581679 | 4.0  | 5.0  | 5.0  | 4.0  | 1.0  | 5.0  | 2.0  | 5.0  | 3.0  | 1.0  |
| 1 | 99540  | 33    | 3.581679 | 5.0  | 5.0  | 5.0  | 4.0  | 5.0  | 3.0  | 4.0  | 4.0  | 3.0  | 5.0  |
| 2 | 99865  | 33    | 3.581679 | 5.0  | 5.0  | 4.0  | 5.0  | 3.0  | 5.0  | 4.0  | 4.0  | 5.0  | 4.0  |
| 3 | 101620 | 33    | 3.581679 | 2.0  | 3.0  | 5.0  | 5.0  | 4.0  | 4.0  | 3.0  | 3.0  | 4.0  | 5.0  |
| 4 | 112974 | 33    | 3.581679 | 5.0  | 5.0  | 5.0  | 5.0  | 5.0  | 3.0  | 5.0  | 5.0  | 5.0  | 3.0  |

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
  - sur1, sur2, sur3, sur4, sur5 ( top 5 similar users who rated that movie.. )
- **Similar movies rated by this user:**
  - smr1, smr2, smr3, smr4, smr5 ( top 5 similar movies rated by this movie.. )
- **UAvg** : User's Average rating
- **MAvg** : Average rating of this movie
- **rating** : Rating of this movie by this user.

Double-click (or enter) to edit

### 4.3.1.2 Featurizing test data

```
# get users, movies and ratings from the Sampled Test
sample_test_users, sample_test_movies, sample_test_ratings = sparse.find(sample_test_sparse_m

sample_train_averages['global']

 3.5902997718873504

start = datetime.now()

if os.path.isfile('sample/small/reg_test.csv'):
    print("It is already created...")
else:

    print('preparing {} tuples for the dataset..\n'.format(len(sample_test_ratings)))
    with open('reg_test.csv', mode='w') as reg_data_file:
        count = 0
        for (user, movie, rating) in zip(sample_test_users, sample_test_movies, sample_test_
            st = datetime.now()

        #----- Ratings of "movie" by similar users of "user" -----
        #print(user, movie)
        try:
            # compute the similar Users of the "user"
            user_sim = cosine_similarity(sample_train_sparse_matrix[user], sample_train_s
            top_sim_users = user_sim.argsort()[:-1][1:] # we are ignoring 'The User' fro
            # get the ratings of most similar users for this movie
            top_ratings = sample_train_sparse_matrix[top_sim_users, movie].toarray().ravel()
            # we will make it's length "5" by adding movie averages to .
            top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5])
            top_sim_users_ratings.extend([sample_train_averages['movie'][movie]]*(5 - len(
            # print(top_sim_users_ratings, end="--")

        except (IndexError, KeyError):
            # It is a new User or new Movie or there are no ratings for given user for to
            ##### Cold Start Problem #####
            top_sim_users_ratings.extend([sample_train_averages['global']]*(5 - len(top_s
            #print(top_sim_users_ratings)
        except:
            print(user, movie)
            # we just want KeyErrors to be resolved. Not every Exception...
            raise

        #----- Ratings by "user" to similar movies of "movie" -----
        try:
            # compute the similar movies of the "movie"
            movie_sim = cosine_similarity(sample_train_sparse_matrix[:, movie].T, sample_t
```

```

top_sim_movies = movie_sim.argsort()[:-1][1:] # we are ignoring 'The User' f
# get the ratings of most similar movie rated by this user..
top_ratings = sample_train_sparse_matrix[user, top_sim_movies].toarray().ravel
# we will make it's length "5" by adding user averages to.
top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:5])
top_sim_movies_ratings.extend([sample_train_averages['user'][user]]*(5-len(to
#rint(top_sim_movies_ratings)
except (IndexError, KeyError):
    #print(top_sim_movies_ratings, end=" : -- ")
    top_sim_movies_ratings.extend([sample_train_averages['global']]*(5-len(top_si
    #print(top_sim_movies_ratings)
except :
    raise

#-----prepare the row to be stores in a file-----
row = list()
# add usser and movie name first
row.append(user)
row.append(movie)
row.append(sample_train_averages['global']) # first feature
#print(row)
# next 5 features are similar_users "movie" ratings
row.extend(top_sim_users_ratings)
#print(row)
# next 5 features are "user" ratings for similar_movies
row.extend(top_sim_movies_ratings)
#print(row)
# Avg_user rating
try:
    row.append(sample_train_averages['user'][user])
except KeyError:
    row.append(sample_train_averages['global'])
except:
    raise
#print(row)
# Avg_movie rating
try:
    row.append(sample_train_averages['movie'][movie])
except KeyError:
    row.append(sample_train_averages['global'])
except:
    raise
#print(row)
# finalley, The actual Rating of this user-movie pair...
row.append(rating)
#print(row)
count = count + 1

# add rows to the file opened..
reg_data_file.write(','.join(map(str, row)))
#print(','.join(map(str, row)))
reg_data_file.write('\n')

```

```

        if (count)%1000 == 0:
            #print(','.join(map(str, row)))
            print("Done for {} rows---- {} ".format(count, datetime.now() - start))
print("",datetime.now() - start)

```

👤 preparing 19273 tuples for the dataset..

```

Done for 1000 rows---- 0:06:00.987447
Done for 2000 rows---- 0:11:55.919001
Done for 3000 rows---- 0:17:52.561195
Done for 4000 rows---- 0:24:10.973360
Done for 5000 rows---- 0:30:49.703708
Done for 6000 rows---- 0:36:44.711092
Done for 7000 rows---- 0:42:44.003191
Done for 8000 rows---- 0:48:38.521152
Done for 9000 rows---- 0:54:37.354014
Done for 10000 rows---- 1:01:49.472084
Done for 11000 rows---- 1:15:41.665944
Done for 12000 rows---- 1:23:30.680912
Done for 13000 rows---- 1:29:39.138605
Done for 14000 rows---- 1:35:48.933873
Done for 15000 rows---- 1:41:58.393987
Done for 16000 rows---- 1:48:15.213625
Done for 17000 rows---- 1:54:49.353169
Done for 18000 rows---- 2:01:18.488841
Done for 19000 rows---- 2:08:36.842630
2:12:25.046441

```

## ▼ Reading from the file to make a test dataframe

```

reg_test_df = pd.read_csv('reg_test.csv', names = ['user', 'movie', 'GAvg', 'sur1', 'sur2', 'sur3', 'sur4', 'sur5', 'smr1', 'smr2', 'smr3', 'smr4', 'smr5', 'UAvg', 'MAvg', 'rating'], header=N
reg_test_df.head(4)

```

|   | user    | movie | GAvg   | sur1   | sur2   | sur3   | sur4   | sur5   | smr1   | smr2   | smr3   | smr4   | smr5   |
|---|---------|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|
| 0 | 808635  | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 |
| 1 | 898730  | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 |
| 2 | 941866  | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 |
| 3 | 1280761 | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 |

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
  - sur1, sur2, sur3, sur4, sur5 ( top 5 similiar users who rated that movie.. )

- **Similar movies rated by this user:**
    - smr1, smr2, smr3, smr4, smr5 ( top 5 simiular movies rated by this movie.. )
  - **UAvg** : User AVerage rating
  - **MAvg** : Average rating of this movie
  - **rating** : Rating of this movie by this user.
- 

Double-click (or enter) to edit

### 4.3.2 Transforming data for Surprise models

```
from surprise import Reader, Dataset
```

#### 4.3.2.1 Transforming train data

- We can't give raw data (movie, user, rating) to train the model in Surprise library.
- They have a saperate format for TRAIN and TEST data, which will be useful for training the mod Surprise.
- We can form the trainset from a file, or from a Pandas DataFrame. <http://surprise.readthedocs.io/en/latest/dataframe.html>

```
# It is to specify how to read the dataframe.
# for our dataframe, we don't have to specify anything extra..
reader = Reader(rating_scale=(1,5))

# create the traindata from the dataframe...
train_data = Dataset.load_from_df(reg_train[['user', 'movie', 'rating']], reader)

# build the trainset from traindata.., It is of dataset format from surprise library..
trainset = train_data.build_full_trainset()
```

#### 4.3.2.2 Transforming test data

- Testset is just a list of (user, movie, rating) tuples. (Order in the tuple is impotant)

```
testset = list(zip(reg_test_df.user.values, reg_test_df.movie.values, reg_test_df.rating.values))
testset[:3]
```

 [(808635, 71, 5), (898730, 71, 3), (941866, 71, 4)]

## 4.4 Applying Machine Learning models

Double-click (or enter) to edit

- Global dictionary that stores rmse and mape for all the models....
    - It stores the metrics in a dictionary of dictionaries

**keys** : model names(string)

**value**: dict(**key** : metric, **value** : value )

```
models_evaluation_train = dict()  
models_evaluation_test = dict()  
  
models_evaluation_train, models_evaluation_test
```

Double-click (or enter) to edit

## Utility functions for running regression models

```

algo.fit(x_train, y_train, eval_metric = 'rmse')
print('Done. Time taken : {} \n'.format(datetime.now()-start))
print('Done \n')

# from the trained model, get the predictions....
print('Evaluating the model with TRAIN data...')
start = datetime.now()
y_train_pred = algo.predict(x_train)
# get the rmse and mape of train data...
rmse_train, mape_train = get_error_metrics(y_train.values, y_train_pred)

# store the results in train_results dictionary..
train_results = {'rmse': rmse_train,
                  'mape' : mape_train,
                  'predictions' : y_train_pred}

#####
# get the test data predictions and compute rmse and mape
print('Evaluating Test data')
y_test_pred = algo.predict(x_test)
rmse_test, mape_test = get_error_metrics(y_true=y_test.values, y_pred=y_test_pred)
# store them in our test results dictionary.
test_results = {'rmse': rmse_test,
                  'mape' : mape_test,
                  'predictions':y_test_pred}

if verbose:
    print('\nTEST DATA')
    print('-'*30)
    print('RMSE : ', rmse_test)
    print('MAPE : ', mape_test)

# return these train and test results...
return train_results, test_results

```

## Utility functions for Surprise modes

```

# it is just to makesure that all of our algorithms should produce same results
# everytime they run...

my_seed = 15
random.seed(my_seed)
np.random.seed(my_seed)

#####
# get (actual_list , predicted_list) ratings given list
# of predictions (prediction is a class in Surprise).
#####
def get_ratings(predictions):

```

```

actual = np.array([pred.r_ui for pred in predictions])
pred = np.array([pred.est for pred in predictions])

return actual, pred

#####
# get 'rmse' and 'mape' , given list of prediction objecs
#####
def get_errors(predictions, print_them=False):

    actual, pred = get_ratings(predictions)
    rmse = np.sqrt(np.mean((pred - actual)**2))
    mape = np.mean(np.abs(pred - actual)/actual)

    return rmse, mape*100

#####
# It will return predicted ratings, rmse and mape of both train and test data  #
#####
def run_surprise(algo, trainset, testset, verbose=True):
    """
    It returns two dictionaries, one for train and the other is for test
    Each of them have 3 key-value pairs, which specify 'rmse', 'mape', and 'predicted'
    """

    start = datetime.now()
    # dictionaries that stores metrics for train and test..
    train = dict()
    test = dict()

    # train the algorithm with the trainset
    st = datetime.now()
    print('Training the model...')
    algo.fit(trainset)
    print('Done. time taken : {}'.format(datetime.now()-st))

    # ----- Evaluating train data-----#
    st = datetime.now()
    print('Evaluating the model with train data..')
    # get the train predictions (list of prediction class inside Surprise)
    train_preds = algo.test(trainset.build_testset())
    # get predicted ratings from the train predictions..
    train_actual_ratings, train_pred_ratings = get_ratings(train_preds)
    # get 'rmse' and 'mape' from the train predictions.
    train_rmse, train_mape = get_errors(train_preds)
    print('time taken : {}'.format(datetime.now()-st))

    if verbose:
        print('*'*15)
        print('Train Data')
        print('*'*15)

```

```

print("RMSE : {}\n\nMAPE : {}".format(train_rmse, train_mape))

#store them in the train dictionary
if verbose:
    print('adding train results in the dictionary..')
train['rmse'] = train_rmse
train['mape'] = train_mape
train['predictions'] = train_pred_ratings

#----- Evaluating Test data-----#
st = datetime.now()
print('\nEvaluating for test data...')
# get the predictions( list of prediction classes) of test data
test_preds = algo.test(testset)
# get the predicted ratings from the list of predictions
test_actual_ratings, test_pred_ratings = get_ratings(test_preds)
# get error metrics from the predicted and actual ratings
test_rmse, test_mape = get_errors(test_preds)
print('time taken : {}'.format(datetime.now()-st))

if verbose:
    print('-'*15)
    print('Test Data')
    print('-'*15)
    print("RMSE : {}\n\nMAPE : {}".format(test_rmse, test_mape))
# store them in test dictionary
if verbose:
    print('storing the test results in test dictionary...')
test['rmse'] = test_rmse
test['mape'] = test_mape
test['predictions'] = test_pred_ratings

print('\n'+'-'*45)
print('Total time taken to run this algorithm :', datetime.now() - start)

# return two dictionaries train and test
return train, test

```

Double-click (or enter) to edit

#### 4.4.1 XGBoost with initial 13 features

```

import xgboost as xgb

# prepare Train data
x_train = reg_train.drop(['user','movie','rating'], axis=1)
y_train = reg_train['rating']

# Dropano Test data

```

```
# prepare test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

# initialize Our first XGBoost model...
first_xgb = xgb.XGBRegressor(silent=False, n_jobs=13, random_state=15, n_estimators=100)
train_results, test_results = run_xgboost(first_xgb, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['first_algo'] = train_results
models_evaluation_test['first_algo'] = test_results

xgb.plot_importance(first_xgb)
plt.show()
```



Training the model..

```
C:\anaconda\lib\site-packages\xgboost\core.py:587: FutureWarning: Series.base is depreca
  if getattr(data, 'base', None) is not None and \
C:\anaconda\lib\site-packages\xgboost\core.py:588: FutureWarning: Series.base is depreca
  data.base is not None and isinstance(data, np.ndarray) \
[16:50:20] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Done. Time taken : 0:00:03.675912
```

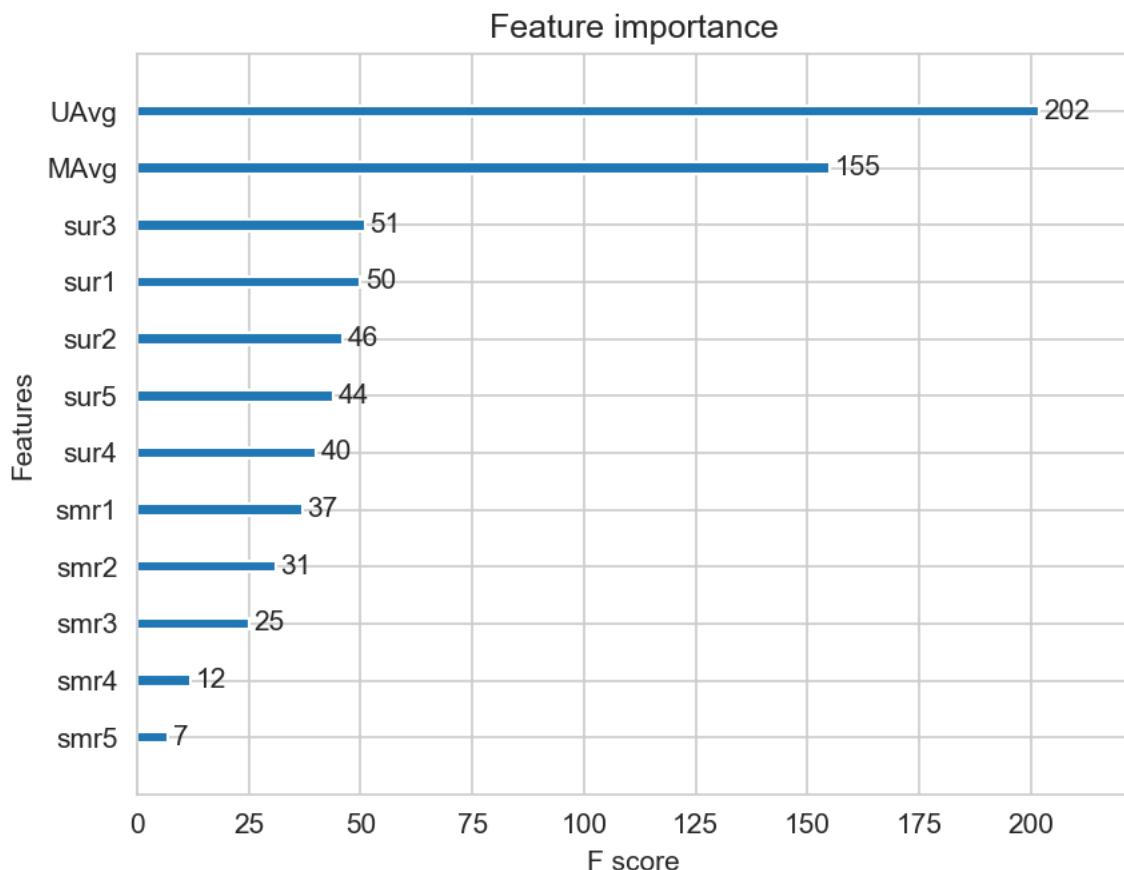
Done

Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA

```
-----  
RMSE : 1.075585551671019  
MAPE : 35.30106791572628
```



## ▼ Hyperparameter Tuning

```
import warnings
warnings.filterwarnings("ignore")
```

```
estimators = [500,1000,1200,1500,2000]
estimators_list = np.asarray(estimators)

for n in estimators_list:

    model = xgb.XGBRegressor(silent=False, n_jobs=-1, random_state=15, n_estimators=n)
    model.fit(x_train,y_train)
    y_train_pred = model.predict(x_train)
    rmse_train, mape_train = get_error_metrics(y_train.values, y_train_pred)
    y_test_pred = model.predict(x_test)
    rmse_test, mape_test = get_error_metrics(y_true=y_test.values, y_pred=y_test_pred)
    print('Best n_estimators == ',n)
    print('Train RMSE : ', rmse_train)
    print('Test RMSE : ', rmse_test)
    print('\n'+ '*'*45)
    print('Train MAPE : ', mape_train)
    print('Test MAPE : ', mape_test)
    print('\n'+ '*'*45)
```



```
[16:58:24] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best n_estimators == 500
Train RMSE : 0.83693905320935
Test RMSE : 1.0758376474932327

-----
Train MAPE : 24.820825684440116
Test MAPE : 34.85146834815715

=====
[16:58:45] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best n_estimators == 1000
Train RMSE : 0.8302001749987319
Test RMSE : 1.0770124055190438

-----
Train MAPE : 24.566256101610705
Test MAPE : 34.68696636170493

=====
[16:59:27] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best n_estimators == 1200
Train RMSE : 0.8267873996309901
Test RMSE : 1.0792939924562814

-----
Train MAPE : 24.433743127719982
Test MAPE : 34.45365790908942

=====
[17:00:17] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best n_estimators == 1500
Train RMSE : 0.8225640787862437
Test RMSE : 1.0876690159542675

-----
Train MAPE : 24.26848426336204
Test MAPE : 33.97430508243837

=====
[17:01:19] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best n_estimators == 2000
Train RMSE : 0.8176241241566515
Test RMSE : 1.084006057501687

-----
Train MAPE : 24.076056450795804
Test MAPE : 34.18315952576197
```

```
import warnings
warnings.filterwarnings("ignore")
```

```
depth = [3,5,7,10,13]
depth_list = np.asarray(depth)

for d in depth_list:

    model = xgb.XGBRegressor(silent=False, n_jobs=-1, random_state=15, n_estimators=1200, max
    model.fit(x_train,y_train)
    y_train_pred = model.predict(x_train)
    rmse_train, mape_train = get_error_metrics(y_train.values, y_train_pred)
    y_test_pred = model.predict(x_test)
    rmse_test, mape_test = get_error_metrics(y_true=y_test.values, y_pred=y_test_pred)
    print('Best depth == ',d)
    print('Train RMSE : ', rmse_train)
    print('Test RMSE : ', rmse_test)
    print('\n'+ '*'*45)
    print('Train MAPE : ', mape_train)
    print('Test MAPE : ', mape_test)
    print('\n'+ '*'*45)
```



```
[17:03:43] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best depth == 3
Train RMSE : 0.8267873996309901
Test RMSE : 1.0792939924562814

-----
Train MAPE : 24.433743127719982
Test MAPE : 34.45365790908942

=====
[17:04:30] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best depth == 5
Train RMSE : 0.7633099006203855
Test RMSE : 1.2116912399328106

-----
Train MAPE : 22.08040545468765
Test MAPE : 31.89347369771194

=====
[17:05:54] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best depth == 7
Train RMSE : 0.6413988790220769
Test RMSE : 1.1719172883632158

-----
Train MAPE : 17.75919135607702
Test MAPE : 32.66317847100031

=====
[17:07:54] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best depth == 10
Train RMSE : 0.33462740341637104
Test RMSE : 1.4051778490749212

-----
Train MAPE : 8.035453707851865
Test MAPE : 36.05172736598685

=====
[17:11:15] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in
Best depth == 13
Train RMSE : 0.06993309369868381
Test RMSE : 1.2026340238757878

-----
Train MAPE : 1.212479226649899
Test MAPE : 32.29704871860572
```

## 4.4.2 Surprise BaselineModel

```
from surprise import BaselineOnly
```

\_\_Predicted\_rating : ( baseline prediction ) \_\_

- [http://surprise.readthedocs.io/en/stable/basic\\_algorithms.html#surprise.prediction\\_algorithms.baseline\\_only.BaselineOnly](http://surprise.readthedocs.io/en/stable/basic_algorithms.html#surprise.prediction_algorithms.baseline_only.BaselineOnly)

$$\hat{r}_{ui} = b_{ui} = \mu + b_u + b_i$$

- $\mu$  : Average of all trainings in training data.
- $b_u$  : User bias
- $b_i$  : Item bias (movie biases)

\_\_Optimization function ( Least Squares Problem ) \_\_

- [http://surprise.readthedocs.io/en/stable/prediction\\_algorithms.html#baselines-estimates-configuration](http://surprise.readthedocs.io/en/stable/prediction_algorithms.html#baselines-estimates-configuration)

$$\sum_{r_{ui} \in R_{train}} (r_{ui} - (\mu + b_u + b_i))^2 + \lambda (b_u^2 + b_i^2) . \text{ [minimize } b_u, b_i \text{ ]}$$

## ▼ Hyperparameter Tuning Surprise BaselineModel

```
from surprise.model_selection import GridSearchCV
import surprise
```

```
start = datetime.now()
```

```
param_grid = {'bsl_options': {'method': ['als']},
              'k': [5, 20, 30, 40, 50, 60],
              'sim_options': {'name': ['pearson_baseline'],
                             'min_support': [2, 3, 4],
                             'shrinkage': [80, 100],
                             'user_based': [True]}}
}
```

```
gs = GridSearchCV (surprise.KNNBaseline, param_grid, measures=['rmse', 'mae'], cv=3, n_jobs=-1)
gs.fit(train_data)
```

```
# best RMSE score
print(gs.best_score['rmse'])
```

```
# combination of parameters that gave the best RMSE score
```

```
print(gs.best_params['rmse'])

print("Time taken = ", datetime.now() - start)

0.9370293747882738
{'bsl_options': {'method': 'als'}, 'k': 60, 'sim_options': {'name': 'pearson_baseline'},
Time taken = 1:06:01.835501

# options are to specify.., how to compute those user and item biases
bsl_options = {'method': 'als',
               'learning_rate': .001,
               'reg': 1
               }
bsl_algo = BaselineOnly(bsl_options=bsl_options)
# run this algorithm.., It will return the train and test results..
bsl_train_results, bsl_test_results = run_surprise(bsl_algo, trainset, testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['bsl_algo'] = bsl_train_results
models_evaluation_test['bsl_algo'] = bsl_test_results

Training the model...
Estimating biases using als...
Done. time taken : 0:00:00.418353

Evaluating the model with train data..
time taken : 0:00:01.709485
-----
Train Data
-----
RMSE : 0.9081297428845365

MAPE : 28.332106114534948

adding train results in the dictionary..

Evaluating for test data...
time taken : 0:00:00.420849
-----
Test Data
-----
RMSE : 1.0791516912122738

MAPE : 35.59112610138221

storing the test results in test dictionary..

-----
Total time taken to run this algorithm : 0:00:02.550842
```

Double-click (or enter) to edit

## 4.4.3 XGBoost with initial 13 features + Surprise Baseline predictor

### Updating Train Data

```
# add our baseline_predicted value as our feature..
reg_train['bslpr'] = models_evaluation_train['bsl_algo']['predictions']
reg_train.head(2)
```

|   | user  | movie | GAvg     | sur1 | sur2 | sur3 | sur4 | sur5 | smr1 | smr2 | smr3 | smr4 | smr5 |
|---|-------|-------|----------|------|------|------|------|------|------|------|------|------|------|
| 0 | 53406 | 33    | 3.581679 | 4.0  | 5.0  | 5.0  | 4.0  | 1.0  | 5.0  | 2.0  | 5.0  | 3.0  | 1.0  |
| 1 | 99540 | 33    | 3.581679 | 5.0  | 5.0  | 5.0  | 4.0  | 5.0  | 3.0  | 4.0  | 4.0  | 3.0  | 5.0  |

### Updating Test Data

```
# add that baseline predicted ratings with Surprise to the test data as well
reg_test_df['bslpr'] = models_evaluation_test['bsl_algo']['predictions']

reg_test_df.head(2)
```

|   | user   | movie | GAvg   | sur1   | sur2   | sur3   | sur4   | sur5   | smr1   | smr2   | smr3   | smr4   | s   |
|---|--------|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|-----|
| 0 | 808635 | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5 |
| 1 | 898730 | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5 |

```
# prepare train data
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# Prepare Test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

# initialize Our first XGBoost model...
xgb_bsl = xgb.XGBRegressor(silent=False, n_jobs=13, random_state=15, n_estimators=100)
train_results, test_results = run_xgboost(xgb_bsl, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_bsl'] = train_results
models_evaluation_test['xgb_bsl'] = test_results

xgb.plot_importance(xgb_bsl)
plt.show()
```



Training the model..

[18:41:48] WARNING: src/objective/regression\_obj.cu:152: reg:linear is now deprecated in Done. Time taken : 0:00:04.883669

Done

Evaluating the model with TRAIN data...

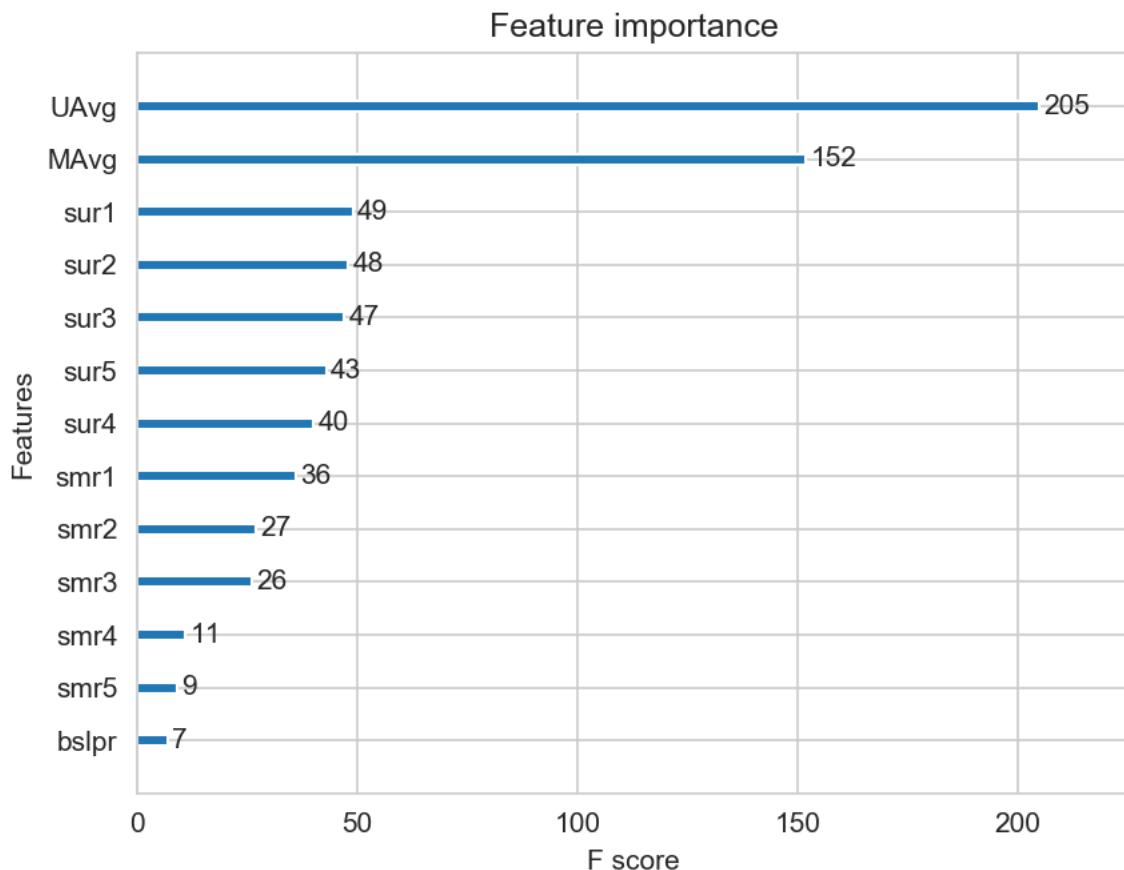
Evaluating Test data

TEST DATA

-----

RMSE : 1.0755536870194393

MAPE : 35.31360443356428



Double-click (or enter) to edit

Double-click (or enter) to edit

#### 4.4.4 Surprise KNNBaseline predictor

```
from surprise import KNNBaseline
```

- KNN BASELINE
  - [http://surprise.readthedocs.io/en/stable/knn\\_inspired.html#surprise.prediction\\_algorithm](http://surprise.readthedocs.io/en/stable/knn_inspired.html#surprise.prediction_algorithm)
- PEARSON\_BASELINE SIMILARITY
  - [http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson\\_b](http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson_b)
- SHRINKAGE
  - 2.2 Neighborhood Models in <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1>
- **predicted Rating : ( \_ based on User-User similarity \_ )**

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot (r_{vi} - b_{vi})}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$

- $b_{ui}$  - Baseline prediction of (user,movie) rating
- $N_i^k(u)$  - Set of **K similar** users (neighbours) of **user (u)** who rated **movie(i)**
- $\text{sim}(u, v)$  - **Similarity** between users **u and v**
  - Generally, it will be cosine similarity or Pearson correlation coefficient.
  - But we use **shrunk Pearson-baseline correlation coefficient**, which is based on the pearsr predictions instead of mean rating of user/item)

Double-click (or enter) to edit

- Predicted rating  ( based on Item Item similarity ):

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in N_u^k(i)} \text{sim}(i, j) \cdot (r_{uj} - b_{uj})}{\sum_{j \in N_u^k(i)} \text{sim}(i, j)}$$

- Notations follows same as above (user user based predicted rating )

#### 4.4.4.1 Surprise KNNBaseline with user user similarities

### ▼ Hyperparameter Tuning KNNBaseline with user user similarities

```
start = datetime.now()
```

```
param_grid = {'bsl_options': {'method': ['als']},
              'k': [5,20, 30,40,50,60],
              'sim_options': {'name': ['pearson_baseline'],
                             'min_support': [2,3,4],
                             'shrinkage':[80,100],
                             'user_based': [True]}}
}

gs = GridSearchCV(surprise.KNNBaseline, param_grid, measures=['rmse', 'mae'], cv=3, n_jobs=-1

gs.fit(train_data)

# best RMSE score
print(gs.best_score['rmse'])

# combination of parameters that gave the best RMSE score
print(gs.best_params['rmse'])

print("Time taken = ", datetime.now() - start)

 0.937377258963861
{'bsl_options': {'method': 'als'}, 'k': 60, 'sim_options': {'name': 'pearson_baseline',
Time taken = 0:57:57.519232

# we specify , how to compute similarities and what to consider with sim_options to our algor
sim_options = {'user_based' : True,
               'name': 'pearson_baseline',
               'shrinkage': 100,
               'min_support': 2
}
# we keep other parameters like regularization parameter and learning_rate as default values.
bsl_options = {'method': 'als'}

knn_bsl_u = KNNBaseline(k=60, sim_options = sim_options, bsl_options = bsl_options)
knn_bsl_u_train_results, knn_bsl_u_test_results = run_surprise(knn_bsl_u, trainset, testset,

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['knn_bsl_u'] = knn_bsl_u_train_results
models_evaluation_test['knn_bsl_u'] = knn_bsl_u_test_results
```



```
Training the model...
Estimating biases using als...
Computing the pearson_baseline similarity matrix...
Done computing similarity matrix.
Done. time taken : 0:00:42.635090

Evaluating the model with train data..
time taken : 0:02:09.167389
-----
Train Data
-----
RMSE : 0.36632663895827805

MAPE : 10.175030164762097

adding train results in the dictionary..

Evaluating for test data...
time taken : 0:00:00.204457
-----
Test Data
-----
RMSE : 1.079140649704147

MAPE : 35.590978793841295

storing the test results in test dictionary..

-----
Total time taken to run this algorithm : 0:02:52.007906
```

#### 4.4.4.2 Surprise KNNBaseline with movie movie similarities

### ▼ Hyperparameter Tuning KNNBaseline with movie movie simila

```
start = datetime.now()

param_grid = {'bsl_options': {'method': ['als']},
              'k': [5,20, 30,40,50,60],
              'sim_options': {'name': ['pearson_baseline'],
                             'min_support': [2,3,4],
                             'shrinkage':[80,100],
                             'user_based': [False]}}
}

gs = GridSearchCV(surprise.KNNBaseline, param_grid, measures=['rmse', 'mae'], cv=3, n_jobs=-1
```

```
gs.fit(train_data)

# best RMSE score
print(gs.best_score['rmse'])

# combination of parameters that gave the best RMSE score
print(gs.best_params['rmse'])

print("Time taken = ", datetime.now() - start)

 0.9818793241277984
{'bsl_options': {'method': 'als'}, 'k': 30, 'sim_options': {'name': 'pearson_baseline',
Time taken = 0:08:45.773554

# we specify , how to compute similarities and what to consider with sim_options to our algor
# 'user_based' : False => this considers the similarities of movies instead of users

sim_options = {'user_based': False,
               'name': 'pearson_baseline',
               'shrinkage': 80,
               'min_support': 3
             }
# we keep other parameters like regularization parameter and learning_rate as default values.
bsl_options = {'method': 'als'}

knn_bsl_m = KNNBaseline(k=30, sim_options = sim_options, bsl_options = bsl_options)

knn_bsl_m_train_results, knn_bsl_m_test_results = run_surprise(knn_bsl_m, trainset, testset,
# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['knn_bsl_m'] = knn_bsl_m_train_results
models_evaluation_test['knn_bsl_m'] = knn_bsl_m_test_results
```



```
Training the model...
Estimating biases using als...
Computing the pearson_baseline similarity matrix...
Done computing similarity matrix.
Done. time taken : 0:00:01.626794

Evaluating the model with train data..
time taken : 0:00:22.693928
-----
Train Data
-----
RMSE : 0.36542560582927563

MAPE : 9.693044100285542

adding train results in the dictionary..

Evaluating for test data...
time taken : 0:00:00.494166
-----
Test Data
-----
RMSE : 1.0792854448413287

MAPE : 35.59336413361821

storing the test results in test dictionary...

-----
Total time taken to run this algorithm : 0:00:24.816881
```

Double-click (or enter) to edit

#### 4.4.5 XGBoost with initial 13 features + Surprise Baseline predictor + KNNBaseline

- - - First we will run XGBoost with predictions from both KNN's ( that uses User\_User and previous features.
    - Then we will run XGBoost with just predictions from both knn models and predictions

##### Preparing Train data

```
# add the predicted values from both knns to this dataframe
reg_train['knn_bsl_u'] = models_evaluation_train['knn_bsl_u']['predictions']
reg_train['knn_bsl_m'] = models_evaluation_train['knn_bsl_m']['predictions']

reg_train.head(2)
```



|   | user  | movie | GAvg     | sur1 | sur2 | sur3 | sur4 | sur5 | smr1 | smr2 | smr3 | smr4 | smr5 |
|---|-------|-------|----------|------|------|------|------|------|------|------|------|------|------|
| 0 | 53406 | 33    | 3.581679 | 4.0  | 5.0  | 5.0  | 4.0  | 1.0  | 5.0  | 2.0  | 5.0  | 3.0  | 1.0  |
| 1 | 99540 | 33    | 3.581679 | 5.0  | 5.0  | 5.0  | 4.0  | 5.0  | 3.0  | 4.0  | 4.0  | 3.0  | 5.0  |

## \_\_Preparing Test data \_\_

```
reg_test_df['knn_bsl_u'] = models_evaluation_test['knn_bsl_u']['predictions']
reg_test_df['knn_bsl_m'] = models_evaluation_test['knn_bsl_m']['predictions']
```

```
reg_test_df.head(2)
```



|   | user   | movie | GAvg   | sur1   | sur2   | sur3   | sur4   | sur5   | smr1   | smr2   | smr3   | s   |
|---|--------|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|-----|
| 0 | 808635 | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5 |
| 1 | 898730 | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5 |

```
# prepare the train data....
```

```
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']
```

```
# prepare the train data....
```

```
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']
```

```
# declare the model
```

```
xgb_knn_bsl = xgb.XGBRegressor(n_jobs=10, random_state=15)
train_results, test_results = run_xgboost(xgb_knn_bsl, x_train, y_train, x_test, y_test)
```

```
# store the results in models_evaluations dictionaries
```

```
models_evaluation_train['xgb_knn_bsl'] = train_results
models_evaluation_test['xgb_knn_bsl'] = test_results
```

```
xgb.plot_importance(xgb_knn_bsl)
plt.show()
```



Training the model..

[20:13:31] WARNING: src/objective/regression\_obj.cu:152: reg:linear is now deprecated in Done. Time taken : 0:00:05.687658

Done

Evaluating the model with TRAIN data...

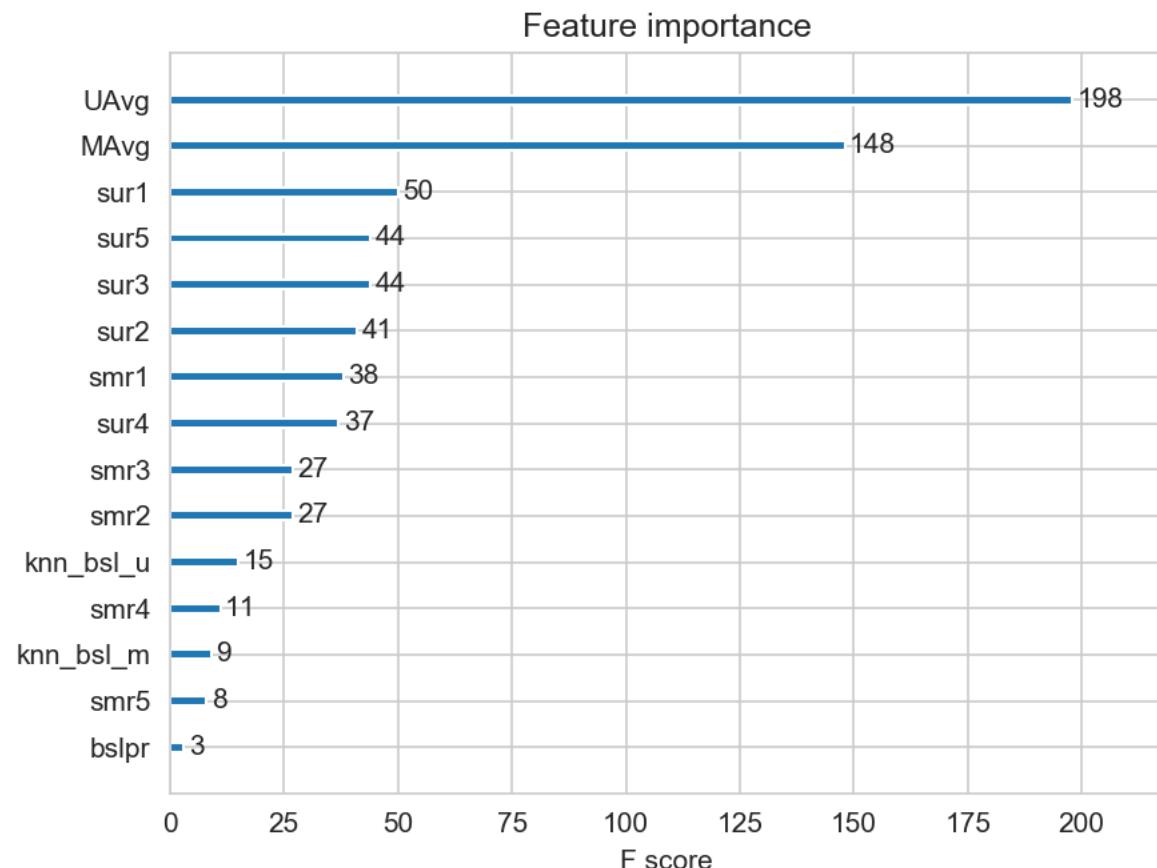
Evaluating Test data

TEST DATA

-----

RMSE : 1.0755585607653444

MAPE : 35.297320345605605



## 4.4.6 Matrix Factorization Techniques

### 4.4.6.1 SVD Matrix Factorization User Movie interactions

```
from surprise import SVD
```

[http://surprise.readthedocs.io/en/stable/matrix\\_factorization.html#surprise.prediction\\_algorithms.m](http://surprise.readthedocs.io/en/stable/matrix_factorization.html#surprise.prediction_algorithms.m)

- \_\_ Predicted Rating : \_\_

- 

- $$\hat{r}_{ui} = \mu + b_u + b_i + q_i^T p_u$$

- $q_i$  - Representation of item(movie) in latent factor space
    - $p_u$  - Representation of user in new latent factor space

- A BASIC MATRIX FACTORIZATION MODEL in <https://datajobs.com/data-science-repo/Recomm>

- Optimization problem with user item interactions and regularization (to avoid overfitting)

- 

- $$\sum_{r_{ui} \in R_{train}} (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2)$$

## ▼ Hyperparameter Tuning SVD

```
from surprise.model_selection import GridSearchCV

# Use movielens-100K

param_grid = {'n_epochs': [5, 10, 15, 20], 'lr_all': [0.002, 0.005],
              'reg_all': [0.4, 0.6], 'n_factors': [60, 80, 100, 120]}
gs = GridSearchCV(SVD, param_grid, measures=['rmse', 'mae'], cv=3)

gs.fit(train_data)

# best RMSE score
print(gs.best_score['rmse'])

# combination of parameters that gave the best RMSE score
print(gs.best_params['rmse'])

👤 0.9512746852704558
{'n_epochs': 20, 'lr_all': 0.005, 'reg_all': 0.4, 'n_factors': 60}

# initialize the model
svd = SVD(n_factors=60, biased=True, random_state=15, verbose=True, n_epochs=20)
svd_train_results, svd_test_results = run_surprise(svd, trainset, testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['svd'] = svd_train_results
models_evaluation_test['svd'] = svd_test_results
```



```
Training the model...
Processing epoch 0
Processing epoch 1
Processing epoch 2
Processing epoch 3
Processing epoch 4
Processing epoch 5
Processing epoch 6
Processing epoch 7
Processing epoch 8
Processing epoch 9
Processing epoch 10
Processing epoch 11
Processing epoch 12
Processing epoch 13
Processing epoch 14
Processing epoch 15
Processing epoch 16
Processing epoch 17
Processing epoch 18
Processing epoch 19
Done. time taken : 0:00:10.698635
```

```
Evaluating the model with train data..
time taken : 0:00:02.169384
```

```
-----
Train Data
```

```
-----
RMSE : 0.7153427264198212
```

```
MAPE : 21.551336861578978
```

```
adding train results in the dictionary..
```

```
Evaluating for test data..
time taken : 0:00:00.511767
```

```
-----
Test Data
```

```
-----
RMSE : 1.0791076662953745
```

```
MAPE : 35.58145698795639
```

```
storing the test results in test dictionary..
```

```
-----
Total time taken to run this algorithm : 0:00:13.379786
```

Double-click (or enter) to edit

#### 4.4.6.2 SVD Matrix Factorization with implicit feedback from user ( user rated movies )

```
from surprise import SVDpp
```

- ----> 2.5 Implicit Feedback in <http://courses.ischool.berkeley.edu/j290-dm/s11/SECURE/a1-kore>
- \_\_ Predicted Rating : \_\_
  - - $\hat{r}_{ui} = \mu + b_u + b_i + q_i^T \left( p_u + |I_u|^{-\frac{1}{2}} \sum_{j \in I_u} y_j \right)$
- $I_u$  -- the set of all items rated by user u
- $y_j$  -- Our new set of item factors that capture implicit ratings.
- **Optimization problem with user item interactions and regularization (to avoid overfitting)**
  - - $\sum_{r_{ui} \in R_{train}} (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2 + \|y_j\|^2)$

## ▼ Hyperparameter Tuning SVD++

```
param_grid = {'n_epochs': [10, 15, 20], 'lr_all': [0.002, 0.005],
             'n_factors': [60, 80, 100, 120]}
gs = GridSearchCV(SVDpp, param_grid, measures=['rmse', 'mae'], cv=3)

gs.fit(train_data)

# best RMSE score
print(gs.best_score['rmse'])

# combination of parameters that gave the best RMSE score
print(gs.best_params['rmse'])

 0.9238133863598583
{'n_epochs': 15, 'lr_all': 0.005, 'n_factors': 60}

# initialize the model
svdpp = SVDpp(n_factors=60, random_state=15, verbose=True, n_epochs=15)
svdpp_train_results, svdpp_test_results = run_surprise(svdpp, trainset, testset, verbose=True

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['svdpp'] = svdpp_train_results
models_evaluation_test['svdpp'] = svdpp_test_results
```

Training the model...

```
processing epoch 0
processing epoch 1
processing epoch 2
processing epoch 3
processing epoch 4
processing epoch 5
processing epoch 6
processing epoch 7
processing epoch 8
processing epoch 9
processing epoch 10
processing epoch 11
processing epoch 12
processing epoch 13
processing epoch 14
```

Done. time taken : 0:01:45.522021

Evaluating the model with train data..

time taken : 0:00:05.868136

-----  
Train Data

-----  
RMSE : 0.6773273263484626

MAPE : 19.96321424247726

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.195476

-----  
Test Data

-----  
RMSE : 1.0795363362492403

MAPE : 35.57547871819873

storing the test results in test dictionary...

-----  
Total time taken to run this algorithm : 0:01:51.585633

Double-click (or enter) to edit

Double-click (or enter) to edit

## 4.4.7 XgBoost with 13 features + Surprise Baseline + Surprise KNNbaseline + MF

### Preparing Train data

```
# add the predicted values from both knns to this dataframe
reg_train['svd'] = models_evaluation_train['svd']['predictions']
reg_train['svdpp'] = models_evaluation_train['svdpp']['predictions']

reg_train.head(2)
```

|   | user  | movie | GAvg     | sur1 | sur2 | sur3 | sur4 | sur5 | smr1 | smr2 | ... | smr4 | smr5 |
|---|-------|-------|----------|------|------|------|------|------|------|------|-----|------|------|
| 0 | 53406 | 33    | 3.581679 | 4.0  | 5.0  | 5.0  | 4.0  | 1.0  | 5.0  | 2.0  | ... | 3.0  | 1.0  |
| 1 | 99540 | 33    | 3.581679 | 5.0  | 5.0  | 5.0  | 4.0  | 5.0  | 3.0  | 4.0  | ... | 3.0  | 5.0  |

2 rows × 21 columns

## ▼ Preparing Test data

```
reg_test_df['svd'] = models_evaluation_test['svd']['predictions']
reg_test_df['svdpp'] = models_evaluation_test['svdpp']['predictions']

reg_test_df.head(2)
```

|   | user   | movie | GAvg   | sur1   | sur2   | sur3   | sur4   | sur5   | smr1   | smr2   | ... | smr4   |
|---|--------|-------|--------|--------|--------|--------|--------|--------|--------|--------|-----|--------|
| 0 | 808635 | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | ... | 3.5903 |
| 1 | 898730 | 71    | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | 3.5903 | ... | 3.5903 |

2 rows × 21 columns

Double-click (or enter) to edit

```
# prepare x_train and y_train
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# prepare test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']
```

```
xgb_final = xgb.XGBRegressor(n_jobs=10, random_state=15)
train_results, test_results = run_xgboost(xgb_final, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_final'] = train_results
```

```
models_evaluation_test['xgb_final'] = test_results
```

```
xgb.plot_importance(xgb_final)
plt.show()
```

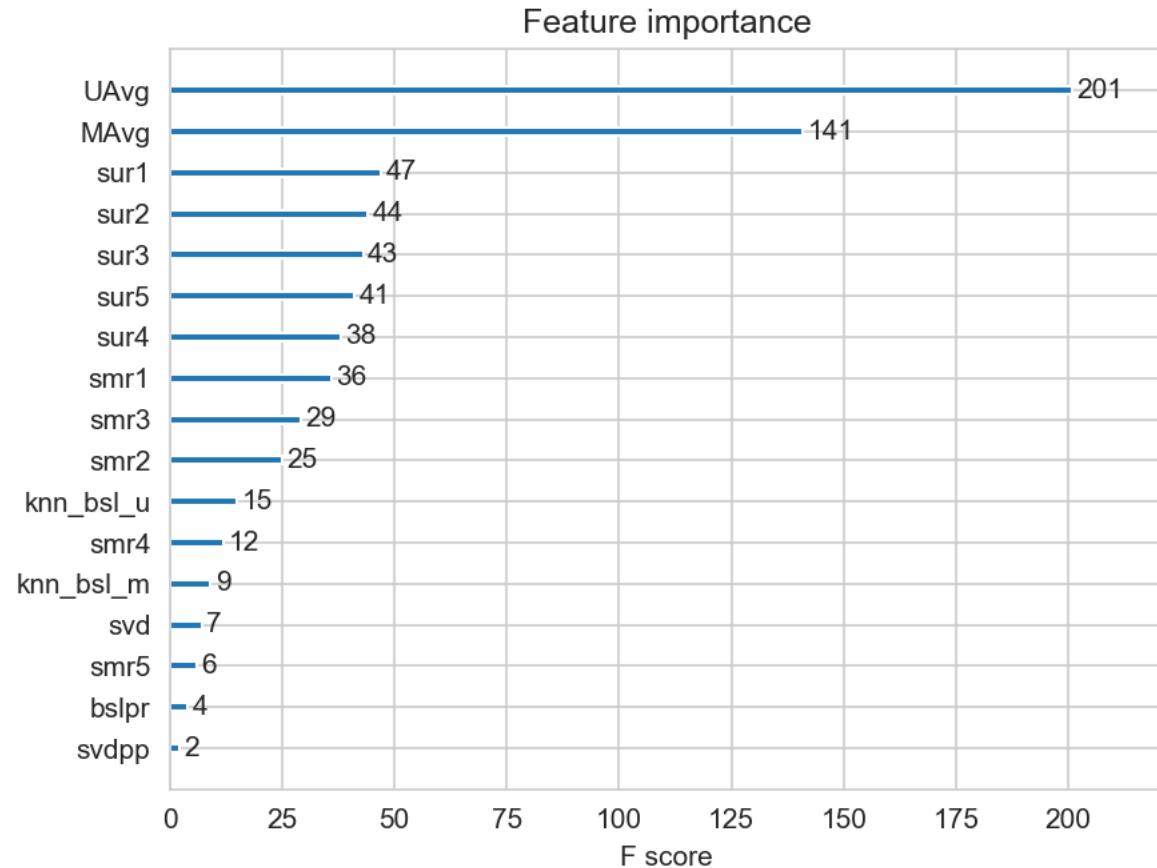
Training the model..  
[22:06:51] WARNING: src/objective/regression\_obj.cu:152: reg:linear is now deprecated in  
Done. Time taken : 0:00:03.744893

Done

Evaluating the model with TRAIN data...  
Evaluating Test data

TEST DATA

RMSE : 1.0754507880473443  
MAPE : 35.28565103694869



#### 4.4.8 XgBoost with Surprise Baseline + Surprise KNNbaseline + MF Techniques

```
# prepare train data
x_train = reg_train[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
y_train = reg_train['rating']
```

```
# test data
x_test = reg_test_df[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
y_test = reg_test_df['rating']

xgb_all_models = xgb.XGBRegressor(n_jobs=10, random_state=15)
train_results, test_results = run_xgboost(xgb_all_models, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_all_models'] = train_results
models_evaluation_test['xgb_all_models'] = test_results

xgb.plot_importance(xgb_all_models)
plt.show()
```



```
Training the model..
```

```
[22:07:10] WARNING: src/objective/regression obj.cu:152: reg:linear is now deprecated in
```

Double-click (or enter) to edit

```
done
```

Double-click (or enter) to edit

```
evaluating test data
```

## 4.5 Comparision between all models

```
IN[5]: 1.0002/04/10014222
```

```
# Saving our TEST_RESULTS into a dataframe so that you don't have to run it again
pd.DataFrame(models_evaluation_test).to_csv('small_sample_results.csv')
models = pd.read_csv('small_sample_results.csv', index_col=0)
models.loc['rmse'].sort_values()
```

 xgb\_final 1.0754507880473443  
xgb\_bsl 1.0755536870194393  
xgb\_knn\_bsl 1.0755585607653444  
first\_algo 1.075585551671019  
svd 1.0791076662953745  
knn\_bsl\_u 1.079140649704147  
bsl\_algo 1.0791516912122738  
knn\_bsl\_m 1.0792854448413287  
svdpp 1.0795363362492403  
xgb\_all\_models 1.0832784518814222  
Name: rmse, dtype: object

Double-click (or enter) to edit

```
globalstart = datetime.now()
print("-"*100)
print("Total time taken to run this entire notebook ( with saved files) is :",datetime.now() -
```



```
-----  
Total time taken to run this entire notebook ( with saved files) is : 0:00:00.000998
```

